



# Cloud Deployment Model Cookbook for Amazon EKS

/ ForgeRock Identity Platform 6.5

Latest update: 6.5.2

Gina Cariaga  
David Goldsmith  
Shankar Raman

ForgeRock AS.  
201 Mission St., Suite 2900  
San Francisco, CA 94105, USA  
+1 415-599-1100 (US)  
[www.forgerock.com](http://www.forgerock.com)

---

Copyright © 2018 ForgeRock AS.

## Abstract

# Step-by-step instructions for getting the CDM up and running on Amazon Elastic Cloud Services for Kubernetes (Amazon EKS).



This work is licensed under the Creative Commons Attribution-NonCommercial-NoDerivs 3.0 Unported License.

To view a copy of this license, visit <https://creativecommons.org/licenses/by-nc-nd/3.0/> or send a letter to Creative Commons, 444 Castro Street, Suite 900, Mountain View, California, 94041, USA.

ForgeRock® and ForgeRock Identity Platform™ are trademarks of ForgeRock Inc. or its subsidiaries in the U.S. and in other countries. Trademarks are the property of their respective owners.

UNLESS OTHERWISE MUTUALLY AGREED BY THE PARTIES IN WRITING, LICENSOR OFFERS THE WORK AS-IS AND MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND CONCERNING THE WORK, EXPRESS, IMPLIED, STATUTORY OR OTHERWISE, INCLUDING, WITHOUT LIMITATION, WARRANTIES OF TITLE, MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NONINFRINGEMENT, OR THE ABSENCE OF LATENT OR OTHER DEFECTS, ACCURACY, OR THE PRESENCE OF ABSENCE OF ERRORS, WHETHER OR NOT DISCOVERABLE. SOME JURISDICTIONS DO NOT ALLOW THE EXCLUSION OF IMPLIED WARRANTIES, SO SUCH EXCLUSION MAY NOT APPLY TO YOU.

EXCEPT TO THE EXTENT REQUIRED BY APPLICABLE LAW, IN NO EVENT WILL LICENSOR BE LIABLE TO YOU ON ANY LEGAL THEORY FOR ANY SPECIAL, INCIDENTAL, CONSEQUENTIAL, PUNITIVE OR EXEMPLARY DAMAGES ARISING OUT OF THIS LICENSE OR THE USE OF THE WORK, EVEN IF LICENSOR HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

DejaVu Fonts

Bitstream Vera Fonts Copyright

Copyright (c) 2003 by Bitstream, Inc. All Rights Reserved. Bitstream Vera is a trademark of Bitstream, Inc.

Permission is hereby granted, free of charge, to any person obtaining a copy of the fonts accompanying this license ("Fonts") and associated documentation files (the "Font Software"), to reproduce and distribute the Font Software, including without limitation the rights to use, copy, merge, publish, distribute, and/or sell copies of the Font Software, and to permit persons to whom the Font Software is furnished to do so, subject to the following conditions:

The above copyright and trademark notices and this permission notice shall be included in all copies of one or more of the Font Software typefaces.

The Font Software may be modified, altered, or added to, and in particular the designs of glyphs or characters in the Fonts may be modified and additional glyphs or characters may be added to the Fonts, only if the fonts are renamed to names not containing either the words "Bitstream" or the word "Vera".

This License becomes null and void to the extent applicable to Fonts or Font Software that has been modified and is distributed under the "Bitstream Vera" names.

The Font Software may be sold as part of a larger software package but no copy of one or more of the Font Software typefaces may be sold by itself.

THE FONT SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT OF COPYRIGHT, PATENT, TRADEMARK, OR OTHER RIGHT. IN NO EVENT SHALL BITSTREAM OR THE GNOME FOUNDATION BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, INCLUDING ANY GENERAL, SPECIAL, INDIRECT, INCIDENTAL, OR CONSEQUENTIAL DAMAGES, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF THE USE OR INABILITY TO USE THE FONT SOFTWARE OR FROM OTHER DEALINGS IN THE FONT SOFTWARE.

Except as contained in this notice, the names of Gnome, the Gnome Foundation, and Bitstream Inc., shall not be used in advertising or otherwise to promote the sale, use or other dealings in this Font Software without prior written authorization from the Gnome Foundation or Bitstream Inc., respectively. For further information, contact: [fonts at gnome dot org](mailto:fonts at gnome dot org).

Arev Fonts Copyright

Copyright (c) 2006 by Tavmjong Bah. All Rights Reserved.

Permission is hereby granted, free of charge, to any person obtaining a copy of the fonts accompanying this license ("Fonts") and associated documentation files (the "Font Software"), to reproduce and distribute the modifications to the Bitstream Vera Font Software, including without limitation the rights to use, copy, merge, publish, distribute, and/or sell copies of the Font Software, and to permit persons to whom the Font Software is furnished to do so, subject to the following conditions:

The above copyright and trademark notices and this permission notice shall be included in all copies of one or more of the Font Software typefaces.

The Font Software may be modified, altered, or added to, and in particular the designs of glyphs or characters in the Fonts may be modified and additional glyphs or characters may be added to the Fonts, only if the fonts are renamed to names not containing either the words "Tavmjong Bah" or the word "Arev".

This License becomes null and void to the extent applicable to Fonts or Font Software that has been modified and is distributed under the "Tavmjong Bah Arev" names.

The Font Software may be sold as part of a larger software package but no copy of one or more of the Font Software typefaces may be sold by itself.

THE FONT SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT OF COPYRIGHT, PATENT, TRADEMARK, OR OTHER RIGHT. IN NO EVENT SHALL TAVMJONG BAH BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, INCLUDING ANY GENERAL, SPECIAL, INDIRECT, INCIDENTAL, OR CONSEQUENTIAL DAMAGES, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF THE USE OR INABILITY TO USE THE FONT SOFTWARE OR FROM OTHER DEALINGS IN THE FONT SOFTWARE.

Except as contained in this notice, the name of Tavmjong Bah shall not be used in advertising or otherwise to promote the sale, use or other dealings in this Font Software without prior written authorization from Tavmjong Bah. For further information, contact: [tavmjong @ free . fr](mailto:tavmjong @ free . fr).

FontAwesome Copyright

Copyright (c) 2017 by Dave Gandy, <http://fontawesome.io>.

This Font Software is licensed under the SIL Open Font License, Version 1.1. See <https://opensource.org/licenses/OFL-1.1>.

# Table of Contents

Preface .....	iv
1. About the ForgeRock Cloud Deployment Model .....	1
1.1. How the CDM Relates to Your Deployment .....	1
1.2. CDM Overview .....	2
1.3. Third-Party Software Deployed With the CDM .....	5
1.4. Best Practices for Implementing the CDM .....	5
2. Setting Up the Deployment Environment .....	7
2.1. An Overview of CDM in Amazon EKS .....	7
2.2. Installing Required Third-Party Software .....	9
2.3. Setting Up an Amazon EKS Environment for CDM .....	9
2.4. Creating and Setting Up a Kubernetes Cluster .....	17
3. Deploying the CDM .....	31
4. Using the CDM .....	47
4.1. Accessing ForgeRock Identity Platform Services .....	47
4.2. Monitoring the CDM .....	50
5. Benchmarking the CDM Performance .....	53
5.1. About CDM Benchmarking .....	53
5.2. Before You Begin .....	55
5.3. Running Directory Services Benchmark Tests .....	59
5.4. Running AM, IDM, or IG Benchmark Tests .....	61
6. Taking the Next Steps .....	82
A. Getting Support .....	83
A.1. ForgeRock DevOps Support .....	83
A.2. Accessing Documentation Online .....	84
A.3. How to Report Problems or Provide Feedback .....	85
A.4. Getting Support and Contacting ForgeRock .....	86
Glossary .....	87

# Preface

The ForgeRock Cloud Deployment Model (CDM) demonstrates a common use ForgeRock Identity Platform™ architecture installed in a DevOps environment. This guide describes the CDM and its default behaviors, and provides steps for replicating the model on Amazon EKS.

For information about how to customize the CDM after you've deployed it, and how to maintain the deployment, see the [Site Reliability Guide for Amazon EKS](#).

## Before You Begin

Before deploying the ForgeRock Identity Platform in a DevOps environment, read the important information in [Start Here](#).

## About ForgeRock Identity Platform Software

ForgeRock Identity Platform™ serves as the basis for our simple and comprehensive Identity and Access Management solution. We help our customers deepen their relationships with their customers, and improve the productivity and connectivity of their employees and partners. For more information about ForgeRock and about the platform, see <https://www.forgerock.com>.

The platform includes the following components:

- ForgeRock® Access Management (AM)
- ForgeRock® Identity Management (IDM)
- ForgeRock® Directory Services (DS)
- ForgeRock® Identity Gateway (IG)

## Chapter 1

# About the ForgeRock Cloud Deployment Model

The `forgeops` repository on GitHub contains artifacts you can use to deploy ForgeRock Identity Platform in a cloud environment. The ForgeRock Cloud Deployment Team has developed Helm charts, Docker images, and other artifacts expressly to build the Cloud Deployment Model (CDM). The CDM is a common use ForgeRock Identity Platform architecture, designed to be easy to deploy and easy to replicate.

This chapter explains how the CDM relates to your deployment, and provides an overview of CDM components and architecture.

## 1.1. How the CDM Relates to Your Deployment

Using the CDM artifacts and *CDM Cookbook* instructions, you can quickly get the ForgeRock Identity Platform running in a Kubernetes cloud environment. The CDM is not a deployment template. But you can deploy the CDM to validate your pre-production deployment against the benchmark results reported in this guide. Then you can customize or scale your deployment environment to meet your specific business requirements. From there, you can develop your own custom template, and write scripts to automate your production deployment.

**Standardizes the process.** The ForgeRock Cloud Deployment Team develops CDM artifacts based on interactions with ForgeRock customers. The Team is made up of technical consultants and cloud software developers who encounter common deployment issues at customer sites. Our mission is to standardize a process for deploying ForgeRock Identity Platform natively in the cloud. We began by standardizing on Kubernetes for our cloud platform.

**Simplifies deployment.** We then developed artifacts such as Helm charts and configuration files to simplify the deployment process. We deployed small-sized, medium-sized, and large-sized production-quality Kubernetes clusters. We conduct continuous integration and continuous deployment as we add new capabilities and fix problems in the system. We maintain, troubleshoot, and tune the system for optimized performance. Most importantly, we have documented the process, and captured benchmark results—a process with results you can replicate.

**Eliminates guesswork.** If you use our CDM artifacts and follow the *CDM Cookbook* instructions without deviation, you can attain results similar to the benchmark results reported in this document. The CDM takes the guesswork out of setting up your cloud environment. It bypasses the deploy-test-integrate-test-repeat cycle many customers struggle through when spinning up the ForgeRock Identity Platform in the cloud for the first time.

## 1.2. CDM Overview

Once you deploy the CDM, the ForgeRock Identity Platform is fully operational within a Kubernetes cluster. Values in the CDM Helm charts override some of the values in the `forgeops` Helm charts. When you deploy CDM, `forgeops` artifacts are overlaid with JVM tuning, memory and CPU limits, and other CDM configurations. Here are some of the benefits of deploying the CDM in your pre-production environment:

### Multi-zone Kubernetes cluster

ForgeRock Identity Platform is deployed in a Kubernetes cluster. You specify one of three cluster sizes:

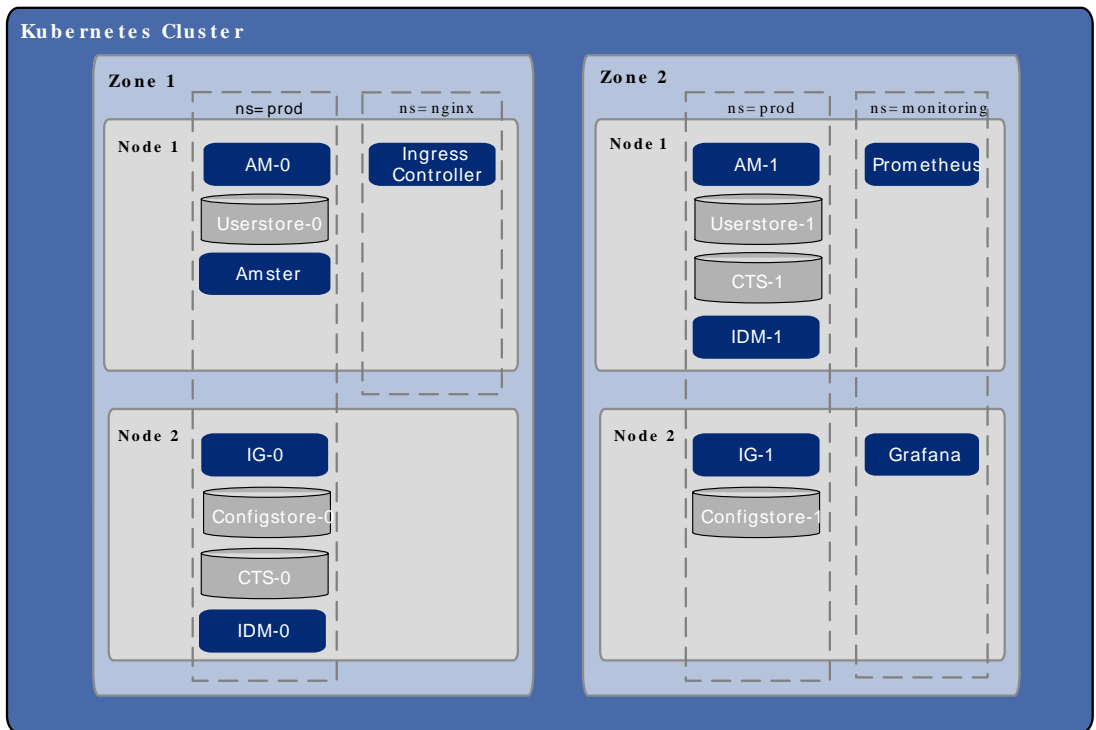
- A small cluster with capacity to handle 1,000,000 test users
- A medium cluster with capacity to handle 10,000,000 test users
- A large cluster with capacity to handle 100,000,000 test users

Two nodes are deployed in the primary zone, and two in an additional zone. This multi-zone, multi-node configuration facilitates scaling and high availability within the cluster.

### Ready-to-use ForgeRock Identity Platform components

- DS instances are deployed with redundancy. Each instance uses separate data stores for users, configuration, and session tokens.
- The `amster` pod facilitates AM deployment by:
  - Obtaining the AM post-installation configuration.
  - Installing the AM server based on configuration from Helm charts in the `forgeops` repository.
  - Importing the post-installation configuration into the AM configuration store.
- AM instances are deployed with redundancy and configured to access the DS data stores.
- IDM instances are deployed with redundancy and configured to work with the DS identity store.
- IG instances are deployed with redundancy.

## CDM Kubernetes Architecture



### Secured communication

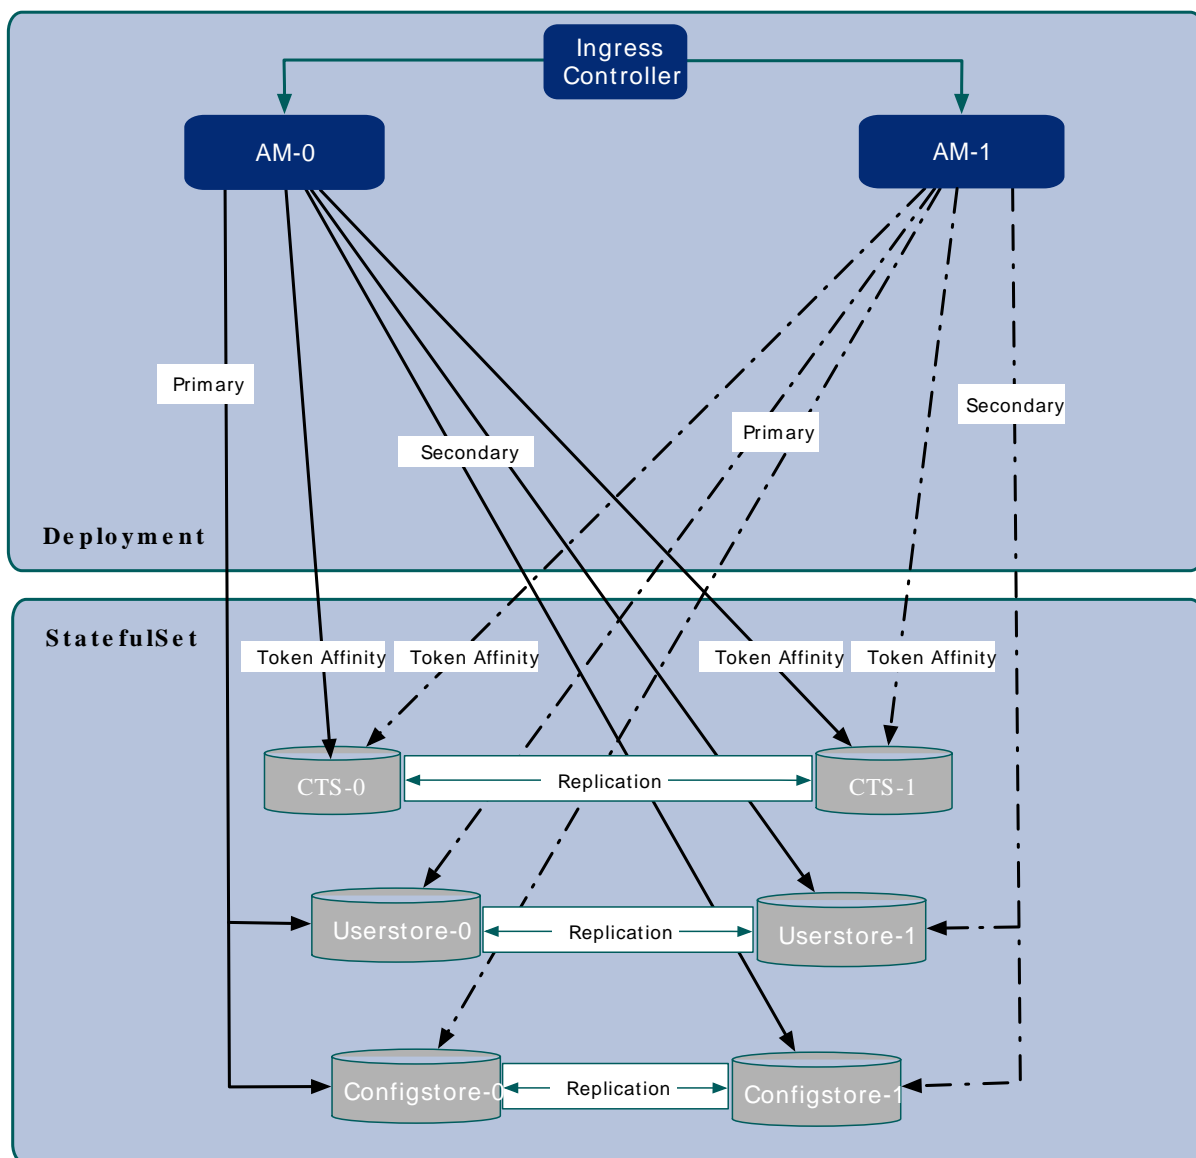
The ingress controller is SSL-enabled. SSL is terminated at the ingress controller. Incoming requests and outgoing responses are encrypted. For more information, see:

- *"Securing Your Deployment"* in the *Site Reliability Guide for GKE*
- *"Securing Your Deployment"* in the *Site Reliability Guide for Amazon EKS*

### StatefulSets

The CDM uses Kubernetes stateful sets to manage the CDM pods. Stateful sets protect against data loss if Kubernetes client containers fail. Primary communication takes place between AM and configuration stores. Secondary communication takes place among AM, configuration stores, and userstores. The CTS data stores are configured for [affinity](#) load balancing for optimal performance. The AM configuration, policies, and application data reside in the configuration store.

## Process Flows in the CDM



## Authentication and authorization

ForgeRock Access Management authentication and OAuth 2 authorization are fully enabled.



## DS replication

All DS instances are configured for full replication of identities, configuration data, and session tokens.

## Backup and restore

The CDM is ready to back up directory data, but backups are not scheduled by default. To schedule backups, see:

- *"Backing Up and Restoring Directory Data"* in the *Site Reliability Guide for GKE*
- *"Backing Up and Restoring Directory Data"* in the *Site Reliability Guide for Amazon EKS*

# 1.3. Third-Party Software Deployed With the CDM

Before you can deploy the CDM, you must install a set of third-party software on your local computer. "Installing Required Third-Party Software" in the *DevOps Release Notes* lists the software you need.

When you deploy the CDM, you deploy not only the ForgeRock Identity Platform in your cluster, but also the following third-party software:

Software	Purpose	More Information
Helm (tiller)	Helm chart deployment	<a href="https://helm.sh">https://helm.sh</a>
Prometheus	Monitoring and Alerting Toolkit	<a href="https://prometheus.io/">https://prometheus.io/</a>
Grafana	Monitoring Metrics Visualization	<a href="https://grafana.com/">https://grafana.com/</a>
Gatling	Load Testing for Benchmarking	<a href="https://gatling.io/">https://gatling.io/</a>
Certificate Manager	Certificate Management	<a href="https://docs.cert-manager.io">https://docs.cert-manager.io</a>

# 1.4. Best Practices for Implementing the CDM

As you work through the *CDM Cookbook* instructions, you can leverage some of the best practices the Cloud Deployment Team has developed over time.

## 1.4.1. Begin at the Beginning

Using the *CDM Cookbook* simplifies deploying ForgeRock Identity Platform in the cloud. But if you deviate from the sequence of steps or customize any of the artifacts, you may not attain the same results the Cloud Deployment Team documented.

### 1.4.2. Provide End-to-End Project Management

Engage at least one deployment professional to provide oversight and guidance during the entire deployment process. We found that at many customer sites, one person architects a deployment plan, then hands off implementation specifications to other team members. Your chances of using the CDM to your best advantage are much higher if everyone on the team understands the context and goals of the CDM.

### 1.4.3. Engage a ForgeRock Cloud Expert

If you intend to implement the CDM as the starting point for a production-quality deployment, then you are probably already working with a ForgeRock technical consultant or partner. If you're not working with a qualified ForgeRock cloud professional, contact your ForgeRock salesperson for more information.

## Chapter 2

# Setting Up the Deployment Environment

This chapter describes how to set up your local computer, how to configure an Amazon EKS environment, and how to create an Amazon EKS cluster before you install the CDM.

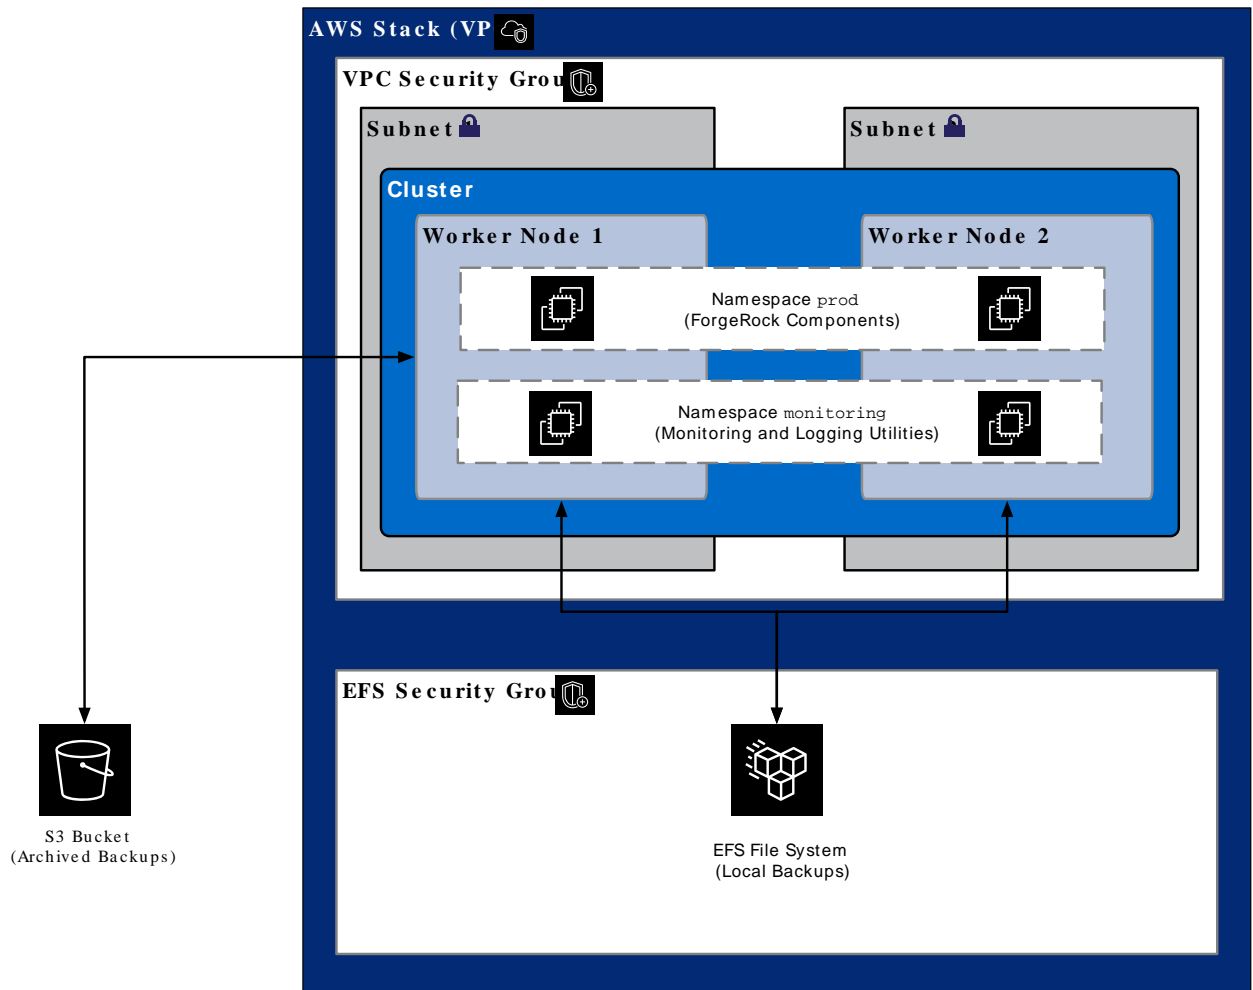
The chapter covers the following topics:

- "An Overview of CDM in Amazon EKS"
- "Installing Required Third-Party Software"
- "Setting Up an Amazon EKS Environment for CDM"
- "Creating and Setting Up a Kubernetes Cluster"

## 2.1. An Overview of CDM in Amazon EKS

The following diagram provides an overview of the CDM deployment in Amazon EKS environment:

## CDM Deployment in Amazon EKS Overview



- An AWS stack template is used to create a virtual private cloud (VPC).
- Two subnets are configured across two availability zones.
- A Kubernetes cluster is created over the two subnets.
- Two worker nodes are created within the cluster. The worker nodes contain the computing infrastructure to run the CDM components.

- An Amazon EFS file system is shared across the worker nodes. The Amazon EFS file system is mounted locally by each DS pod for storing directory data backup.
- An Amazon S3 bucket is used to archive the directory backup data.

## 2.2. Installing Required Third-Party Software

Before installing the CDM, you must obtain non-ForgeRock software and install it on your local computer.

The CDM has been validated with specific third-party software versions. Review "[Installing Required Third-Party Software](#)" in the *DevOps Release Notes* to determine which software you need. Then install the software on your local computer.

The CDM *might* also work with older or newer versions of the third-party software.

## 2.3. Setting Up an Amazon EKS Environment for CDM

The CDM runs in a Kubernetes cluster in an Amazon EKS environment.

This section outlines the steps that the Cloud Deployment Team performed to set up an Amazon EKS environment for deploying CDM:

- "Cloning the forgeops Repository"
- "Granting Permissions to Configure CDM Resources"
- "Creating a Virtual Private Cloud"
- "Creating an Amazon EFS File System"
- "Creating an Amazon S3 Bucket Policy"
- "Creating a Service Role"
- "Creating a Key Pair to Connect to Worker Nodes"

### 2.3.1. Cloning the forgeops Repository

Before you can deploy the CDM environment, you must clone the `forgeops` repository. The `samples` directory in the `forgeops` repository contains the CDM configuration.

#### *To Obtain the forgeops Repository*

The `forgeops` repository is a public Git repository. You do not need credentials to clone it:

1. Clone the `forgeops` repository:

```
$ git clone https://github.com/ForgeRock/forgeops.git
```

2. Check out the `release/6.5.2` branch:

```
$ cd forgeops  
$ git checkout release/6.5.2
```

### 2.3.2. Granting Permissions to Configure CDM Resources

To deploy and manage CDM, you need to create resources such as a Virtual Private Cloud (VPC), Elastic Filestore (EFS), S3 bucket, and so on.

This section outlines how the Cloud Deployment Team granted permissions enabling users to manage CDM resources in Amazon EKS.

#### *To Grant Users AWS Permissions*

1. Create a group with the name `cdm-group`.
2. Attach the following AWS preconfigured policies to the `cdm-group` group:
  - `AWSLambdaFullAccess`
  - `IAMUserChangePassword`
  - `IAMReadOnlyAccess`
  - `PowerUserAccess`
3. Create the following four policies in the IAM service of your AWS account:
  - a. Create the `eks-full-access` policy using the `eks-full-access.json` file in the `/path/to/forgeops/etc/aws-example-iam-policies` directory.
  - b. Create the `iam-change-user-key` policy using the `iam-change-user-key.json` file in the `/path/to/forgeops/etc/aws-example-iam-policies` directory.
  - c. Create the `iam-create-role` policy using the `iam-create-role.json` file in the `/path/to/forgeops/etc/aws-example-iam-policies` directory.
  - d. Create the `iam-limited-write` policy using the `iam-limited-write.json` file in the `/path/to/forgeops/etc/aws-example-iam-policies` directory.
4. Attach the policies you created to the `cdm-group` group.
5. Assign AWS users who will set up CDM to the `cdm-group` group.

#### Note

Perform all the steps in this chapter as an AWS user who is a member of the `cdm-group` group.

### 2.3.3. Creating a Virtual Private Cloud

This section outlines how the Cloud Deployment Team created a Virtual Private Cloud (VPC) for deploying CDM in an Amazon EKS environment.

#### *To Create a Virtual Private Cloud*

1. In a terminal window on your local machine, set up your **aws** command-line interface environment using the **aws configure** command.
2. Get the user's group membership and verify that the user is a member of the `cdm-group` AWS group:

```
$ aws iam list-groups-for-user --user-name my-sre.user --output json
{
  "Groups": [
    {
      "Path": "/",
      "CreateDate": "2019-03-11T21:03:17Z",
      "GroupId": "ABCDEFGHIJKLMN0PQRST",
      "Arn": "arn:aws:iam::123456789012:group/cdm-group",
      "GroupName": "cdm-group"
    }
  ]
}
```

3. View the network configurations in the `amazon-eks-vpc.yaml` VPC template YAML file in the `/path/to/forgeops/etc` directory. Edit the parameters to suit your requirements.
4. Set your default region to `us-east-1`:

```
$ aws configure set default.region us-east-1
```

The Cloud Deployment Team deployed the CDM in the `us-east-1` (N. Virginia) region. Use this region when deploying the CDM, regardless of your actual location, if you want to validate your deployment against the benchmarks in "*Benchmarking the CDM Performance*".

To use any other region, note the following:

- The region must support Amazon EKS.
- Objects required for your EKS cluster must reside in the same region. To make sure that the objects are created in the correct region, be sure to set your default region as shown above.
- The AMI ID in the `NodeImageId` parameter must be modified to use an equivalent AMI ID in your region in the **aws cloudformation deploy** command to create worker nodes.

5. In a terminal window, change to the directory where your `forgeops` repository clone resides, and then run the `aws cloudformation deploy` command:

```
$ cd /path/to/forgeops
$ aws cloudformation deploy \
  --stack-name my-clust-vpc \
  --template-file ./etc/amazon-eks-vpc.yaml \
  --capabilities CAPABILITY_IAM
Waiting for changeset to be created.
Waiting for stack create/update to complete
Successfully created/updated stack - my-clust-vpc
```

6. You will need the following VPC parameter values when you create a cluster:

- **SecurityGroups:** This allows the Amazon EKS cluster to communicate with your worker nodes. This security group is also referred to as `control-plane-security-group`.
- **VpcId:** This is used for creating the subnets in your Amazon Elastic Compute Cloud (Amazon EC2) account.
- **SubnetIds:** This is used to create subnets into which the worker nodes are launched.

To obtain the values, run the following command:

```
$ aws cloudformation describe-stacks \
  --stack-name my-clust-vpc \
  --output table --query 'Stacks[0].Outputs'
...
OutputKey      |  OutputValue
-----
SecurityGroups |  sg-0fc026c50480dd635
VpcId           |  vpc-0b94e8afa7984aa15
SubnetIds       |  subnet-0f7f430272343f04a,subnet-056b1cf2f6038db88
...
```

You can write your own script to automate VPC creation. For an example script, see `eks-create-vpc.sh`.

### 2.3.4. Creating an Amazon EFS File System

Amazon EKS relies on the Amazon Elastic File System (Amazon EFS) for storing directory data backups. When you create an Amazon EFS file system, the mount points are created on the two zones in your region, to make the file system highly available.

#### *To Create an Amazon EFS File System*

1. Create an EFS security group:



```
$ aws ec2 create-security-group \
--description "Security group used for NFS mount" \
--group-name my-eks-efs-sg \
--vpc-id your-vpc-id \
{
  "GroupId": "sg-0b701714496115e94"
}
```

2. Note the **GroupId** value output from the EFS security group creation command. You will need it when you grant access to worker nodes, and when you create the Kubernetes cluster.
3. Create an Amazon EFS file system:

```
$ aws efs create-file-system --performance-mode maxIO \
--creation-token my-eks-nfs-mount1
{
  "SizeInBytes": {
    "Value": 0
  },
  "ThroughputMode": "bursting",
  "CreationToken": "my-eks-nfs-mount1",
  "Encrypted": false,
  "CreationTime": 1542135708.0,
  "PerformanceMode": "maxIO",
  "FileSystemId": "fs-dcdb9c96",
  "NumberOfMountTargets": 0,
  "LifeCycleState": "creating",
  "OwnerId": "890123456789"
}
```

Note the **FileSystemId** value. You need it to configure the directory server backup.

4. Create a mount point to the Amazon EFS file system in each of your subnets, and note the **MountTargetId** values from each mount point:

```
$ aws efs create-mount-target \
--file-system-id your-FileSystemId \
--subnet-id subnet-id-1 \
{
  "MountTargetId": "fsm-t-87ad7bcc",
  "NetworkInterfaceId": "eni-03599eae2aa18127",
  "FileSystemId": "fs-dcdb9c96",
  "LifeCycleState": "creating",
  "SubnetId": "subnet-0f7f430272343f04a",
  "OwnerId": "890123456789",
  "IpAddress": "192.168.81.44"
}
```

```
$ aws efs create-mount-target \
  --file-system-id your-FileSystemId \
  --subnet-id subnet-id-2
{
  "MountTargetId": "fsmt-b8ad7bf3",
  "NetworkInterfaceId": "eni-0b40712f7bbc9959b",
  "FileSystemId": "fs-dcdb9c96",
  "LifeCycleState": "creating",
  "SubnetId": "subnet-056b1cf2f6038db88",
  "OwnerId": "890123456789",
  "IpAddress": "122.168.162.96"
}
```

5. Grant permissions for worker nodes to access the EFS file system:

```
$ aws ec2 authorize-security-group-ingress \
  --group-id efs-security-group-id \
  --protocol tcp --port 2049 \
  --source-group control-plane-security-group
```

You can write your own script to automate Amazon EFS configuration. For an example script, see [eks-create-filesystem.sh](#).

### 2.3.5. Creating an Amazon S3 Bucket Policy

An Amazon Simple Storage Service (Amazon S3) bucket policy lets you manage access to the objects in the bucket.

#### *To Create an Amazon S3 Bucket Policy*

1. Create an Amazon S3 bucket:

```
$ aws s3 mb s3://myfops1
```

2. Create the `my-s3policy.json` JSON file with the following text:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "s3:*",
      "Resource": [
        "arn:aws:s3:::myfops1",
        "arn:aws:s3:::myfops1/*"
      ]
    }
  ]
}
```

3. Create the Amazon S3 policy:

```
$ aws iam create-policy \
  --policy-name my-FR-Sync-Policy \
  --policy-document file://my-s3policy.json
{
  "Policy": {
    "PolicyName": "my-FR-Sync-Policy",
    "PermissionsBoundaryUsageCount": 0,
    "CreateDate": "2018-11-13T19:23:00Z",
    "AttachmentCount": 0,
    "IsAttachable": true,
    "PolicyId": "ANPAI55MF7UP4NHRIZUXE",
    "DefaultVersionId": "v1",
    "Path": "/",
    "Arn": "arn:aws:iam::813759318741:policy/my-FR-Sync-Policy",
    "UpdateDate": "2018-11-13T19:23:00Z"
  }
}
```

Note the **Arn** value. It is required to create worker nodes.

4. Update the bucket permissions to block public access:

```
$ aws s3api put-public-access-block \
  --public-access-block-configuration BlockPublicAcls=true,IgnorePublicAcls=true,\
  BlockPublicPolicy=true,RestrictPublicBuckets=true \
  --bucket myfops1
```

You can write your own script to automate Amazon S3 policy configuration. For an example script, see [eks-create-s3-policy.sh](#).

### 2.3.6. Creating a Service Role

A role that a service assumes to perform actions on your behalf is called a service role. The CDM uses a service role named **eksServiceRole**.

#### *To Create a Service Role*

1. Make sure you've checked out the release/6.5.2 branch of the **forgeops** repository.
2. In the terminal window, change to the **etc** subdirectory in your **forgeops** repository clone:

```
$ cd /path/to/forgeops/etc
```

3. Using a text editor, create the **eks-service-role.json** JSON file with the following content:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": "sts:AssumeRole",
      "Effect": "Allow",
      "Principal": {
        "Service": "eks.amazonaws.com"
      }
    }
  ]
}
```

4. Create the `eksServiceRole` service role by running the following **aws** command:

```
$ aws iam create-role \
--role-name eksServiceRole \
--description "Allows Amazon EKS to manage clusters on your behalf." \
--assume-role-policy-document file://eks-service-role.json
{
  "Role": {
    "AssumeRolePolicyDocument": {
      "Version": "2012-10-17",
      "Statement": [
        {
          "Action": "sts:AssumeRole",
          "Effect": "Allow",
          "Principal": {
            "Service": "eks.amazonaws.com"
          }
        }
      ]
    },
    "RoleId": "AR0AJGZJTA76K3Z0ZSC5G",
    "CreateDate": "2018-11-13T20:05:44Z",
    "RoleName": "eksServiceRole",
    "Path": "/",
    "Arn": "arn:aws:iam::813759318741:role/eksServiceRole"
  }
}
```

Note the service role **Arn** value. You need this value to create a cluster.

5. Attach the Amazon Cluster and Service policies to the role by running the following two **aws** commands:

```
$ aws iam attach-role-policy \
--role-name eksServiceRole \
--policy-arn arn:aws:iam::aws:policy/AmazonEKSClusterPolicy
```

```
$ aws iam attach-role-policy \
--role-name eksServiceRole \
--policy-arn arn:aws:iam::aws:policy/AmazonEKSServicePolicy
```

You can write your own script to automate service role configuration. For an example script, see `eks-create-service-role.sh`.

### 2.3.7. Creating a Key Pair to Connect to Worker Nodes

You need a key pair to connect to the worker nodes that you create for CDM in an Amazon EKS environment. AWS does not store the private key, so you should download the private key when you create the key pair.

#### *To Create a Key Pair to Connect to Amazon Worker Nodes*

1. Create a key pair on AWS and download the corresponding PEM file locally:

```
$ aws ec2 create-key-pair --key-name my-eks-clust-key \
  --query 'KeyMaterial' --output text > my-eks-clust-key.pem
```

This **aws** command creates a key pair in AWS with the name *my-eks-clust-key* and downloads the *my-eks-clust-key.pem* file to the current local directory.

2. Change permissions on the downloaded PEM file to disable access for others.

```
$ chmod 400 my-eks-clust-key.pem
```

You can write your own script to create a key pair and download PEM file locally. For an example script, see [eks-create-keypair.sh](#).

## 2.4. Creating and Setting Up a Kubernetes Cluster

This section describes how to create and set up a Kubernetes cluster that can run the CDM, and covers the following topics:

- "Creating a Kubernetes Cluster"
- "Setting the Kubernetes Context to Your Amazon EKS Cluster"
- "Creating Amazon EKS Worker Nodes"
- "Mounting the EFS Filesystem on Worker Nodes"
- "Creating Namespaces"
- "Configuring RBAC for the Cluster"
- "Deploying an Ingress Controller"
- "Creating Storage Classes"
- "Deploying the Certificate Manager"
- "Deploying Monitoring Infrastructure"

## 2.4.1. Creating a Kubernetes Cluster

This section outlines how the Cloud Deployment Team created our cluster.

### To Create a Kubernetes Cluster for CDM

1. Identify the values that you will need when you run the command to create a Kubernetes cluster in the next step:
  - **Name.** Any string that you want to use as the name of your Kubernetes cluster, for example, *my-eks-cluster*.
  - **Kubernetes Version.** Specify the minor Kubernetes version from the table in "Choosing the Kubernetes Version" in the *DevOps Release Notes*.
  - **Role ARN.** The Amazon Resource Name (ARN) of the security role you created in "To Create a Service Role".
  - **Subnet Ids.** The list of *SubnetIds* that you noted after you created the VPC by running "To Create a Virtual Private Cloud".
  - **Security Group Ids.** The *SecurityGroups* (also referred to as the *Control Plane Security Group*) that you noted after you created the VPC by running "To Create a Virtual Private Cloud".
2. Create the Kubernetes cluster using the **aws eks create-cluster** command:

```
$ aws eks create-cluster --name my-eks-cluster \
--kubernetes-version "1.11" \
--role-arn "arn:aws:iam::813759318741:role/eksServiceRole" \
--resources-vpc-config "subnetIds=your-subnet-ids,securityGroupIds=your-control-plane-security-group-id"
{
  "cluster": {
    "status": "CREATING",
    "name": "my-eks-cluster",
    "certificateAuthority": {},
    "roleArn": "arn:aws:iam::813759318741:role/eksServiceRole",
    "resourcesVpcConfig": {
      "subnetIds": [
        "subnet-0f7f430272343f04a",
        "subnet-056b1cf2f6038db88"
      ],
      "vpcId": "vpc-0b94e8afa7984aa15",
      "securityGroupIds": [
        "sg-0fc026c50480dd635"
      ]
    },
    "version": "1.11",
    "arn": "arn:aws:eks:us-east-1:813759318741:cluster/my-eks-cluster",
    "platformVersion": "eks.2",
    "createdAt": 1542139923.019
  }
}
```

You can write your own script to automate Amazon EKS cluster configuration. For an example script, see [eks-create-cluster.sh](#), which is called by the example environment setup script, [eks-up.sh](#).

## 2.4.2. Setting the Kubernetes Context to Your Amazon EKS Cluster

Set up the Kubernetes context on your local computer to use the Amazon EKS cluster you created, to ensure that your deployments are directed to the correct Kubernetes cluster.

### *To Set the Kubernetes Context to the Amazon EKS Cluster*

1. Verify that you have version 1.16.18 or later of the AWS CLI installed:

```
$ aws --version
aws-cli/1.16.22 Python/2.7.10 Darwin/18.2.0 botocore/1.12.12
```

2. Create or update the Kubernetes context configuration in your local machine:

```
$ aws eks update-kubeconfig --name my-eks-cluster
Added new context arn:aws:eks:us-east-1:123456789012:cluster/my-eks-cluster to /Users/your.user/.kube/config
```

## 2.4.3. Creating Amazon EKS Worker Nodes

Amazon EKS worker nodes are standard Amazon EC2 instances, and run in your Amazon EKS cluster. This section outlines how the Amazon EKS worker nodes are configured for deploying CDM.

### *To Create Amazon EKS Worker Nodes*

1. Identify the values that you will need when you run the command to create Amazon EKS worker nodes in the next step:
  - **Stack Name.** The worker nodes will use a separate stack. So *do not use the same stack you used before*.
  - **Key Name.** The name of the key pair you created in "To Create a Key Pair to Connect to Amazon Worker Nodes".
  - **AMI ID.** The ID of the Amazon EKS-Optimized AMI you will use to create worker nodes. Choose the AMI ID for the latest patch version of Kubernetes 1.11 for your region from the tables in Amazon EKS-Optimized AMI.
  - **Cluster Name.** The name of the cluster you created in "To Create a Kubernetes Cluster for CDM".
  - **VPC ID.** The [VPCId](#) of the VPC you created in "To Create a Virtual Private Cloud".
  - **SubnetId.** The [SubnetIds](#) that you noted after you created the VPC in "To Create a Virtual Private Cloud".

- **ClusterControlPlaneSecurityGroup.** The **SecurityGroup** of your VPC that you noted after you created the VPC in "To Create a Virtual Private Cloud".
- **S3 Policy ARN.** The ARN of the S3 policy that you created in "To Create an Amazon S3 Bucket Policy".
- **EFS Security Group.** The security group (**GroupId**) of the EFS you created in "To Create an Amazon EFS File System".

2. Run the **aws cloudformation deploy** command to create nodes of the same size.

- Use the following example command to create small nodes (1,000,000 users). Make sure you specify **NodeInstanceType="m5.xlarge"**:

```
$ aws cloudformation deploy \
--stack-name my-small-nodes \
--template-file /path/to/forgedops/etc/amazon-eks-nodegroup.yaml \
--parameter-overrides KeyName=my-eks-clust-key \
NodeImageId="my-ami-id" \
NodeInstanceType="m5.xlarge" \
NodeAutoScalingGroupMinSize=2 \
NodeAutoScalingGroupMaxSize=4 \
NodeVolumeSize=50 \
ClusterName="my-eks-cluster" \
NodeGroupName="prod-nodegroup" \
ClusterControlPlaneSecurityGroup=your-cluster-plane-security-group \
VpcId=your-vpc-id \
Subnets=your-vpc-subnet-ids \
S3PolicyArn=your-s3-policy-arn \
EFSSecurityGroup=your-efs-security-group \
--capabilities CAPABILITY_IAM
Waiting for changeset to be created..
Waiting for stack create/update to complete
Successfully created/updated stack - my-small-nodes
```

- Use the following example command to create medium nodes (10,000,000 users). Make sure you specify **NodeInstanceType="m5.4xlarge"**:



```
$ aws cloudformation deploy \
--stack-name my-medium-nodes \
--template-file path/to/forgeops/etc/amazon-eks-nodgroup.yaml \
--parameter-overrides KeyName=my-eks-clust-key \
NodeImageId="my-ami-id" \
NodeInstanceType="m5.4xlarge" \
NodeAutoScalingGroupMinSize=2 \
NodeAutoScalingGroupMaxSize=4 \
NodeVolumeSize=50 \
ClusterName="my-eks-cluster" \
NodeGroupName="prod-nodgroup" \
ClusterControlPlaneSecurityGroup=your-ec2-security-group \
VpcId=your-vpc \
Subnets=your-vpc-subnetids \
S3PolicyArn=your-s3-policy-arn \
EFSSecurityGroup=your-efs-security-group \
--capabilities CAPABILITY_IAM
Waiting for changeset to be created..
Waiting for stack create/update to complete
Successfully created/updated stack - my-medium-nodes
```

- Use the following example command to create large nodes (100,000,000 users). Make sure you specify `NodeInstanceType="c5.9xlarge"`:

```
$ aws cloudformation deploy \
--stack-name my-large-nodes \
--template-file path/to/forgeops/etc/amazon-eks-nodgroup.yaml \
--parameter-overrides KeyName=my-eks-clust-key \
NodeImageId="my-ami-id" \
NodeInstanceType="c5.9xlarge" \
NodeAutoScalingGroupMinSize=2 \
NodeAutoScalingGroupMaxSize=4 \
NodeVolumeSize=50 \
ClusterName="my-eks-cluster" \
NodeGroupName="prod-nodgroup" \
ClusterControlPlaneSecurityGroup=your-ec2-security-group \
VpcId=your-vpc \
Subnets=your-vpc-subnetids \
S3PolicyArn=your-s3-policy-arn \
EFSSecurityGroup=your-efs-security-group \
--capabilities CAPABILITY_IAM
Waiting for changeset to be created..
Waiting for stack create/update to complete
Successfully created/updated stack - my-medium-nodes
```

3. Get the `NodeInstanceRole` ARN for the worker nodes you created.

```
$ aws cloudformation describe-stacks \
--stack-name my-small-nodes \
--query 'Stacks[0].Outputs[0].OutputValue' \
--output text
arn:aws:iam::813759318741:role/my-small-nodes-NodeInstanceRole-1454COW0GD82G
```

4. Create a configuration map to link master and worker nodes.

```
$ kubectl apply -f - <<EOF
apiVersion: v1
kind: ConfigMap
metadata:
  name: aws-auth
  namespace: kube-system
data:
  mapRoles: |
    - rolearn: arn:aws:iam::813759318741:role/my-small-nodes-NodeInstanceRole-1454COW0GD82G
      username: system:node:{{EC2PrivateDNSName}}
      groups:
        - system:bootstrappers
        - system:nodes
EOF
configmap/aws-auth created
```

You can write your own script to automate worker node configuration. For an example script, see [eks-create-worker-nodes.sh](#), which is called by the example environment setup script, [eks-up.sh](#).

For instructions on granting cluster access to additional users, see "Granting Access to Multiple Users" in the *Site Reliability Guide for Amazon EKS*.

## 2.4.4. Mounting the EFS Filesystem on Worker Nodes

To create a mount point and mount the EFS filesystem on each worker node, follow this procedure:

### *To Mount the EFS Filesystem on Worker Nodes*

1. Get the security group ID of the worker nodes:

```
$ aws ec2 describe-security-groups \
  --filters Name=group-name,Values=*your-worker-node-stack-name* \
  --query "SecurityGroups[*].{ID:GroupId}" | grep ID
```

For example:

```
$ aws ec2 describe-security-groups \
  --filters Name=group-name,Values=*my-small-nodes* \
  --query "SecurityGroups[*].{ID:GroupId}" | grep ID
"ID": "sg-0321055d76624b08b"
```

2. Get a list of [MountTargetId](#) values for the EFS filesystem you noted in "To Create an Amazon EFS File System":

```
$ aws efs describe-mount-targets \
  --file-system-id file-system-id | grep MountTargetId
```

3. Add the worker node security group to the security groups parameter on each mount target:

```
$ aws efs modify-mount-target-security-groups \
  --mount-target-id first-MountTargetId \
  --security-groups worker-node-security-group-id efs-security-group-id
```

```
$ aws efs modify-mount-target-security-groups \
  --mount-target-id second-MountTargetId \
  --security-groups worker-node-security-group-id efs-security-group-id
```

4. Enable inbound SSH traffic to the worker nodes:

```
$ aws ec2 authorize-security-group-ingress \
  --group-id worker-node-security-group-id \
  --protocol tcp --port 22 --cidr 0.0.0.0/0 || true
```

5. Get the public IP addresses of the worker nodes using the AWS Console.

6. Get the region name you are using to deploy the CDM:

```
$ aws configure get region
```

7. Mount the EFS filesystem on the first worker node using its public IP address, and create the `tmp` directory. Make sure to specify the `your-key-pair`, `your-worker-node1-ip`, and `your-region` values from your environment:

```
$ ssh -o StrictHostKeyChecking=no -i your-key-pair.pem ec2-user@your-worker-node1-ip /bin/bash <<EOF
sudo mount -t nfs your-filesystemID.efs.your-region.amazonaws.com: /mnt
sudo mkdir -p /mnt/export/tmp
EOF
```

You can write your own script to automate mounting the EFS filesystem on worker nodes. For an example script, see [eks-mount-efs.sh](#), which is called by the example environment setup script, [eks-up.sh](#).

## 2.4.5. Creating Namespaces

The CDM uses the `monitoring` and `prod` namespaces. The `monitoring` namespace contains monitoring and notification tools and related resources. The `prod` namespace contains all the other resources of the CDM.

For more information on Kubernetes namespaces, see [Namespaces](#) in the Kubernetes Concepts documentation.

### To Create Namespaces for CDM

1. Create the `monitoring` namespace:

```
$ kubectl create namespace monitoring
namespace/monitoring created
```

2. Create the `prod` namespace:

```
$ kubectl create namespace prod
namespace/prod created
```

3. Set context to the `prod` namespace:

```
$ kubectl config set-context $(kubectl config current-context) --namespace=prod
```

You can write your own script to automate namespace creation. A section of the `eks-up.sh` script contains example commands.

## 2.4.6. Configuring RBAC for the Cluster

Role-based access control (RBAC) in Kubernetes enables you to control how users access the API resources running on your cluster.

For information about RBAC in Kubernetes, see [Using RBAC Authorization](#) in the Kubernetes documentation.

To replicate the RBAC implemented in CDM, perform the following procedure:

### *To Configure RBAC*

1. Create the `tiller` service account:

```
$ kubectl --namespace kube-system create serviceaccount tiller
serviceaccount/tiller created
```

2. Bind the `tiller` service account to the `cluster-admin` role:

```
$ kubectl create clusterrolebinding tiller --clusterrole cluster-admin \
--serviceaccount=kube-system:tiller
clusterrolebinding.rbac.authorization.k8s.io/tiller created
```

3. Initialize the `tiller` service account:

```
$ helm init --upgrade --wait --service-account tiller
$HELM_HOME has been configured at /Users/.../.helm.
```

Tiller (the Helm server-side component) has been installed into your Kubernetes Cluster.

Please note: by default, Tiller is deployed with an insecure 'allow unauthenticated users' policy. To prevent this, run `helm init` with the `--tiller-tls-verify` flag.

For more information on securing your installation see: [https://docs.helm.sh/using\\_helm/#securing-your-helm-installation](https://docs.helm.sh/using_helm/#securing-your-helm-installation)  
Happy Helming!

You can write your own script to automate RBAC configuration. For an example script, see `helm-rbac-init.sh`, which is called by the example environment setup script, `eks-up.sh`.

## 2.4.7. Deploying an Ingress Controller

An ingress controller must be running in the Kubernetes cluster before the ingress rules are applied and cluster resources can be accessed.

For information about the ingress controller used in CDM, see [NGINX Ingress Controller](#).

The ingress controller can be configured with or without a static IP address. The IP address resides in the same region as the CDM deployment.

When you deploy an ingress controller in production, you can set up an array of CIDRs to allow communication from only a particular IP or range of IP addresses.

For details about how to configure CIDR in your own production deployment, see "Controlling Access by Configuring a CIDR Block" in the *Site Reliability Guide for Amazon EKS*.

## To Deploy the Ingress Controller

1. Install the **nginx** ingress controller using the following **helm install** command:

```
$ helm install --namespace nginx --name nginx \
--set rbac.create=true \
--set controller.publishService.enabled=true \
--set controller.stats.enabled=true \
--set controller.service.externalTrafficPolicy=Local \
--set controller.service.type=LoadBalancer \
--set controller.image.tag="0.21.0" \
--set controller.service.annotations."service\.beta\.kubernetes\.io/aws-load-balancer-type"="nlb" \
stable/nginx-ingress
NAME:      nginx
LAST DEPLOYED: Thu Nov 15 13:52:50 2019
NAMESPACE: nginx
STATUS:    DEPLOYED
...

NOTES:
The nginx-ingress controller has been installed.
It may take a few minutes for the LoadBalancer IP to be available.
You can watch the status by running 'kubectl --namespace nginx get services -o wide -w nginx-nginx-ingress-controller'
...

apiVersion: v1
kind: Secret
metadata:
  name: example-tls
  namespace: foo
data:
  tls.crt: <base64 encoded cert>
  tls.key: <base64 encoded key>
type: kubernetes.io/tls
```

2. Verify that the ingress controller has been installed, and obtain the DNS name of the ingress controller:

```
$ kubectl --namespace nginx get services nginx-nginx-ingress-controller \
--no-headers -o custom-columns=NAME:.status.loadBalancer.ingress[0].hostname
acbe5956ee92011e8b5f91285b865ace-8644dc79ba786353.elb.us-east-1.amazonaws.com
```

If the command returns no value, then it means the controller installation has not yet completed. Repeat the **kubectl** command again after a few minutes, and it will return the DNS name of the ingress controller.

3. Obtain the hosted zone ID for your domain:

```
$ aws route53 list-hosted-zones-by-name --dns-name your-domain \
--query 'HostedZones[0].Id'
```

For example:

```
$ aws route53 list-hosted-zones-by-name --dns-name example.com \
--query 'HostedZones[0].Id'
"/hostedzone/Z3CRFEHJA9N5QU"
```

In the output, the last portion, **Z3CRFEHJA9N5QU**, is the hosted zone ID.

#### 4. Create routing configuration files for AM, IDM, IG, and Gatling components:

##### a. Create the **rt53-template.json** JSON file with the following content:

```
{
  "Comment": "UPSERT a record ",
  "Changes": [
    {
      "Action": "UPSERT",
      "ResourceRecordSet": {
        "Name": "URL-head.prod.Your-domain",
        "Type": "CNAME",
        "TTL": 300,
        "ResourceRecords": [
          {
            "Value": "DNS-name"
          }
        ]
      }
    }
  ]
}
```

##### b. Replace the following values in the **rt53-template.json** file:

- **Your-domain** with your DNS domain name.
- **DNS-name** with the DNS name of the ingress controller you obtained in Step 2 of this procedure.

##### c. Make four copies of the **rt53-template.json** JSON file, with the names **am-route53.json**, **idm-route53.json**, **ig-route53.json**, and **gatling-route53.json**.

```
$ cp rt53-template.json am-route53.json; \
cp rt53-template.json idm-route53.json; \
cp rt53-template.json ig-route53.json; \
cp rt53-template.json gatling-route53.json
```

- Edit the **am-route53.json** file and replace *URL-head* with **login** in the value of the **Name** parameter.
- Edit the **idm-route53.json** file and replace *URL-head* with **openidm** in the value of the **Name** parameter.
- Edit the **ig-route53.json** file and replace *URL-head* with **openig** in the value of the **Name** parameter.

- g. Edit the `gatling-route53.json` file and replace `URL-head` with `gatling` in the value of the `Name` parameter.
5. Create routing records for the components to point to the cluster URL:
  - a. Create a routing record for AM:

```
$ aws route53 change-resource-record-sets \
  --hosted-zone-id your-hosted-zone-id \
  --change-batch file://am-route53.json
{
  "ChangeInfo": {
    "Status": "PENDING",
    "Comment": "UPsert a record ",
    "SubmittedAt": "2018-11-16T21:27:11.269Z",
    "Id": "/change/C1R7BEXLYLT295"
  }
}
```

- b. Create a routing record for IDM:

```
$ aws route53 change-resource-record-sets \
  --hosted-zone-id your-hosted-zone-id \
  --change-batch file://idm-route53.json
```

- c. Create a routing record for IG:

```
$ aws route53 change-resource-record-sets \
  --hosted-zone-id your-hosted-zone-id \
  --change-batch file://ig-route53.json
```

- d. Create a routing record for Gatling:

```
$ aws route53 change-resource-record-sets \
  --hosted-zone-id your-hosted-zone-id \
  --change-batch file://gatling-route53.json
```

You can write your own script to automate ingress controller deployment. For an example script, see `eks-create-ingress-cntlr.sh`, which is called by the example environment setup script, `eks-up.sh`.

## 2.4.8. Creating Storage Classes

Kubernetes storage classes define the characteristics of the storage volumes used in an environment. CDM uses the following storage classes:

- `fast` - used by the `userstore` pod
- `fast10` - used by the `ctsstore` pod
- `standard` - used by the `configstore` pod
- `nfs` - used by the DS pods to store the directory data backup.

For more information on Kubernetes storage classes, see [Storage Classes](#) in the Kubernetes Concepts documentation.

## To Create Storage Classes

1. Create the `fast`, `standard`, and `fast10` storage classes using the following `kubectl` command:

```
$ kubectl create -f - <<EOF
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: fast
provisioner: kubernetes.io/aws-ebs
parameters:
  type: gp2
---
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: standard
provisioner: kubernetes.io/aws-ebs
parameters:
  type: gp2
---
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: fast10
provisioner: kubernetes.io/aws-ebs
parameters:
  type: io1
  fstype: ext4
  iopsPerGB: "10"
---
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: nfs
provisioner: kubernetes.io/aws-efs
EOF
storageclass.storage.k8s.io/fast created
storageclass.storage.k8s.io/standard created
storageclass.storage.k8s.io/fast10 created
storageclass.storage.k8s.io/nfs created
```

2. The `gp2` storage class is created automatically by AWS when you create storage classes. This storage class is not needed for the CDM, so delete it:

```
$ kubectl delete storageclass gp2
```

3. Set the `standard` storage class as the default:

```
$ kubectl patch storageclass standard -p \
'{"metadata": {"annotations":{"storageclass.kubernetes.io/is-default-class":"true"}}}'
```



You can write your own script to automate storage class creation. For an example script, see `eks-create-sc.sh`, which is called by the example environment setup script, `eks-up.sh`.

## 2.4.9. Deploying the Certificate Manager

In the CDM, certificate management is provided by the `cert-manager` add-on. In the Amazon EKS environment, CDM is configured to use the self-signed certificate issued by `cert-manager` to secure communication.

For details about how to secure communication in your own production deployment, see "Automating Certificate Management" in the *Site Reliability Guide for Amazon EKS*.

### *To Deploy the Certificate Manager*

1. Deploy the Kubernetes `cert-manager` add-on:

```
$ helm install stable/cert-manager --namespace kube-system --version v0.5.0
NAME:      brawny-marmot
LAST DEPLOYED: Tue Jul 30 13:46:55 2019
NAMESPACE: kube-system
STATUS:    DEPLOYED
. . .
```

2. Run the `kubectl get pods -n kube-system` command to determine when the certificate manager is available.

Examine the `kubectl get pods` output and look for a pod whose name contains the string `cert-manager`. If this pod has `Running` status, the certificate manager is available. Note that you might have to run the `kubectl get pods` command several times before the `cert-manager` pod attains `Running` status.

## 2.4.10. Deploying Monitoring Infrastructure

The CDM uses Prometheus and Grafana for CDM monitoring and reporting.

### *To Deploy CDM Monitoring Tools*

1. Refresh your local Helm repository cache:

```
$ helm repo update
```

2. Change to the directory that contains configuration values for installing Prometheus:

```
$ cd /path/to/forgeops/etc/prometheus-values
```

3. Install the `prometheus-operator` chart:

```
$ helm install stable/prometheus-operator \
  --name monitoring-prometheus-operator --values prometheus-operator.yaml \
  --set=rbac.install=true --namespace=monitoring
NAME:      monitoring-prometheus-operator
LAST DEPLOYED: Sun Jul 28 06:34:15 2019
NAMESPACE: monitoring
STATUS: DEPLOYED
. . .
```

4. Change to the directory that contains the `forgerock-metrics` Helm chart:

```
$ cd /path/to/forgoeps/helm
```

5. Install the `forgerock-metrics` chart:

```
$ helm install forgerock-metrics \
  --name=monitoring-forgoeps-metrics \
  --set=rbac.install=true --namespace=monitoring
NAME:      invincible-panther
LAST DEPLOYED: Sun Jul 28 06:38:30 2019
NAMESPACE: monitoring
STATUS: DEPLOYED
. . .
```

You can write your own script to automate monitoring infrastructure deployment. For an example script, see `deploy-prometheus.sh`, which is called by the example environment setup script, `eks-up.sh`.

## Chapter 3

# Deploying the CDM

Now that you've set up your deployment environment following the instructions in the previous chapter, you're ready to deploy the CDM. This chapter shows you how to deploy the CDM in your Kubernetes cluster using artifacts from the `forgeops` repository.

Perform the following procedures:

1. "To Customize the Deployment YAML Files"
2. "To Set Your Kubernetes Context"
3. "To Deploy Supporting Components"
4. "To Deploy DS"
5. "To Deploy AM"
6. "To Deploy IDM"
7. "To Deploy IG"
8. "To Deploy the web Application"
9. "To Scale the Deployment"

You can write your own script to automate completing all these procedures. For an example script, see `deploy.sh`.

After successfully deploying the CDM, you can access AM, IDM, and DS using graphical and command-line interfaces. For more information, see *"Using the CDM"*.

When you're finished working with the CDM, you can remove it from your cluster. See *"Deleting ForgeRock Identity Platform From Your Namespace"* in the *DevOps Developer's Guide* for details.

### *To Customize the Deployment YAML Files*

Make the following changes to the Helm charts in your `forgeops` repository clone:

1. Make sure you've checked out the release/6.5.2 branch of the `forgeops` repository.
2. In your clone of the `forgeops` repository, change to the directory containing the appropriate artifacts:

- If you're deploying the small cluster: `/path/to/forgeops/samples/config/prod/s-cluster`
- If you're deploying the medium cluster: `/path/to/forgeops/samples/config/prod/m-cluster`
- If you're deploying the large cluster: `/path/to/forgeops/samples/config/prod/l-cluster`

3. Make the following changes to the `common.yaml` file:

- Change the value of the `domain` key to your own domain, making sure you put a period before the domain name.
- Change the value of the `fqdn` key to `login.prod.domain`, where `domain` is your own domain.

For example:

```
domain: .mydomain.com
fqdn: login.prod.mydomain.com
. . .
```

4. In the `nfs` section of the `dsadmin.yaml` file, change these key values:

- For GKE, set these key values:
  - For `server`, specify the Cloud Filestore's IP address.
  - For `path`, specify `/export`.

In this example, the Cloud Filestore's IP address is `10.198.53.26`:

```
. . .
nfs:
  server: 10.198.53.26
  path: /export
. . .
```

For more information about CDM's Cloud Filestore requirement, see "Setting up a GCP Project for the CDM" in the *Cloud Deployment Model Cookbook for GKE*.

- For Amazon EKS, set these key values:
  - For `server`, specify the file system ID of your Amazon EFS file system.
  - For `path`, specify `/export`.

In this example, the domain qualified file system ID of Amazon EFS storage `fs-dcdb9c96.efs.us-east-1.amazonaws.com`:

```
. . .
nfs:
  server: fs-dcdb9c96.efs.us-east-1.amazonaws.com
  path: /export
. . .
```

For more information about how to set up an Amazon EFS storage for CDM, see "Creating an Amazon EFS File System".

## To Set Your Kubernetes Context

Ensure that your Kubernetes context is set to the Kubernetes cluster and namespace in which you want to deploy the CDM:

1. Run the **kubectx** command to determine the Kubernetes cluster in your current context:

```
$ kubectx
```

- a. Examine the **kubectx** command output. Highlighted text indicates the Kubernetes cluster in your current context.
- b. If the Kubernetes cluster in your current context does not match the cluster you want to deploy to, set the cluster in your current context to the correct cluster. See "Setting up a Kubernetes Context" in the *DevOps Developer's Guide*.

2. Run the **kubens** command to determine the Kubernetes namespace in your current context. For example:

```
$ kubens
default
kube-public
kube-system
monitoring
nginx
prod
```

- a. Examine the **kubens** command output. Highlighted output indicates the Kubernetes namespace set in your current context.
- b. If the Kubernetes namespace in your current context is not the **prod** namespace, run the following command to set the **prod** namespace in your current context:

```
$ kubens prod
```

## To Deploy Supporting Components

1. Make sure you've checked out the release/6.5.2 branch of the **forgeops** repository.
2. In your clone of the **forgeops** repository, change to the directory containing the appropriate artifacts:
  - If you're deploying the small cluster: `/path/to/forgeops/samples/config/prod/s-cluster`
  - If you're deploying the medium cluster: `/path/to/forgeops/samples/config/prod/m-cluster`
  - If you're deploying the large cluster: `/path/to/forgeops/samples/config/prod/l-cluster`

3. Install the `frconfig` chart, which creates Kubernetes objects used by other ForgeRock pods:

```
$ helm install --name prod-frconfig --namespace prod \
  --values common.yaml --values frconfig.yaml /path/to/forgeops/helm/frconfig
NAME:      prod-frconfig
LAST DEPLOYED: Wed Jul 31 16:00:31 2019
NAMESPACE: prod
STATUS:    DEPLOYED

RESOURCES:
==> v1/Secret
NAME      AGE
frconfig  0s

==> v1/ConfigMap
frconfig  0s
```

4. Install the `dsadmin` chart, which creates a persistent volume claim for the NFS storage class, mounts the shared-access file storage service that holds directory data backups, and archives backups to cloud storage:

```
$ helm install --name prod-dsadmin --namespace prod \
  --values common.yaml --values dsadmin.yaml /path/to/forgeops/helm/dsadmin
NAME:      prod-dsadmin
LAST DEPLOYED: Wed Jul 31 16:10:16 2019
NAMESPACE: prod
STATUS:    DEPLOYED

RESOURCES:
==> v1/PersistentVolumeClaim
NAME      AGE
ds-backup 1s

==> v1/Deployment
dsadmin   0s

==> v1beta1/CronJob
gcs-sync  0s

==> v1/Pod(related)

NAME                                READY  STATUS   RESTARTS  AGE
dsadmin-69fb874f74-ncgz4           0/1    Pending  0          0s

==> v1/PersistentVolume
NAME      AGE
ds-backup 1s
```

## To Deploy DS

Install the `ds` chart three times to create DS servers for the configuration, user, and CTS stores:

1. Install the AM configuration store:

```
$ helm install --name prod-configstore --namespace prod \
  --values common.yaml --values configstore.yaml /path/to/forgeops/helm/ds
NAME:      prod-configstore
LAST DEPLOYED: Wed Jul 31 16:13:57 2019
NAMESPACE: prod
STATUS: DEPLOYED

RESOURCES:
==> v1/ConfigMap
NAME      AGE
configstore 0s

==> v1/Service
configstore 0s

==> v1beta1/StatefulSet
configstore 0s

==> v1/Pod(related)

NAME            READY  STATUS   RESTARTS  AGE
configstore-0   0/1    Pending  0          0s

==> v1/Secret

NAME      AGE
configstore 0s
```

## 2. Install the AM user store:

```
$ helm install --name prod-userstore --namespace prod \
  --values common.yaml --values userstore.yaml /path/to/forgeops/helm/ds
NAME:      prod-userstore
LAST DEPLOYED: Wed Jul 31 16:14:56 2019
NAMESPACE: prod
STATUS: DEPLOYED

RESOURCES:
==> v1/Service
NAME      AGE
userstore 0s

==> v1beta1/StatefulSet
userstore 0s

==> v1/Pod(related)

NAME            READY  STATUS   RESTARTS  AGE
userstore-0     0/1    Pending  0          0s

==> v1/Secret

NAME      AGE
userstore 0s

==> v1/ConfigMap
userstore 0s
```

### Note

For information about provisioning the user store, see:

- "Before You Begin" in the *Cloud Deployment Model Cookbook for GKE*
- "Before You Begin"

3. Install the AM CTS store, which holds AM session tokens, OAuth 2.0 tokens, and session blacklists and whitelists:

```
$ helm install --name prod-ctsstore --namespace prod \
  --values common.yaml --values ctsstore.yaml /path/to/forgeops/helm/ds
NAME:      prod-ctsstore
LAST DEPLOYED: Wed Jul 31 16:15:53 2019
NAMESPACE: prod
STATUS:    DEPLOYED

RESOURCES:
==> v1beta1/StatefulSet
NAME          AGE
ctsstore      0s

==> v1/Pod(related)

NAME          READY   STATUS    RESTARTS   AGE
ctsstore-0    0/1     Pending   0           0s

==> v1/Secret
NAME          AGE
ctsstore      0s

==> v1/ConfigMap
ctsstore      0s

==> v1/Service
ctsstore      0s
```

4. Review the logs from the directory server pods to verify that the following DS pods started up successfully:

- configstore-0
- configstore-1
- userstore-0
- userstore-1
- ctsstore-0



- `ctsstore-1`

For example, the following command verifies that the `configstore-0` pod started successfully:

```
$ kubectl logs configstore-0 | grep successfully
[31/Jul/2019:22:50:50 +0000] category=CORE severity=NOTICE msgID=135 msg=The Directory Server has
started successfully
[31/Jul/2019:22:50:50 +0000] category=CORE severity=NOTICE msgID=139 msg=The Directory Server has
sent an alert notification generated by class org.opens.server.core.DirectoryServer (alert type org
.opens.server.DirectoryServerStarted, alert ID org.opens.messages.core-135): The Directory Server
has started successfully
```

This step is complete when you have verified that the pods started successfully.

## To Deploy AM

1. Install the `openam` chart:

```
$ helm install --name prod-openam --namespace prod \
--values common.yaml --values openam.yaml /path/to/forgeops/helm/openam
NAME:      prod-openam
LAST DEPLOYED: Thu Aug  1 12:26:58 2019
NAMESPACE: prod
STATUS:    DEPLOYED

RESOURCES:
==> v1/ConfigMap
NAME          AGE
am-configmap  1s
boot-json     1s

==> v1/Service
openam        1s

==> v1beta1/Deployment
prod-openam-openam  1s

==> v1beta1/Ingress
openam             1s

==> v1/Pod(related)

NAME                                READY  STATUS   RESTARTS  AGE
prod-openam-openam-6f846bbf69-gt82c  0/1    Init:0/3  0         1s

==> v1/Secret

NAME          AGE
openam-secrets  1s
```

2. Install the `amster` chart:

```
$ helm install --name prod-amster --namespace prod \
--values common.yaml --values amster.yaml /path/to/forgeops/helm/amster
NAME:      prod-amster
LAST DEPLOYED: Thu Aug  1 12:28:00 2019
NAMESPACE: prod
STATUS:    DEPLOYED

RESOURCES:
==> v1/Secret
NAME                AGE
amster-secrets      0s

==> v1/ConfigMap
amster-config        0s
amster-prod-amster   0s

==> v1beta1/Deployment
amster               0s

==> v1/Pod(related)

NAME                                READY  STATUS   RESTARTS  AGE
amster-7b8b6d9667-hxq57            0/2    Init:0/1  0          0s
```

### 3. Check the status of the pods in the deployment until all pods are ready:

#### a. Run the **kubectl get pods** command:

```
$ kubectl get pods
NAME                                READY    STATUS    RESTARTS  AGE
amster-7b8b6d9667-hxq57            2/2      Running   0          1m
configstore-0                      1/1      Running   0          10m
configstore-1                      1/1      Running   0          10m
ctsstore-0                         1/1      Running   0          9m
ctsstore-1                        1/1      Running   0          9m
dsadmin-69fb874f74-dln6b           1/1      Running   0          11m
prod-openam-openam-6f846bbf69-gt82c 1/1      Running   0          2m
userstore-0                       1/1      Running   0          10m
userstore-1                       1/1      Running   0          9m
```

#### b. Review the output. Deployment is complete when:

- The **READY** column indicates all containers are available. The entry in the **READY** column represents [total number of containers/number of available containers].
- All entries in the **STATUS** column indicate **Running**.

#### c. If necessary, continue to query your deployment's status until all the pods are ready.

### 4. Review the Amster pod's log to determine whether the deployment completed successfully.

Use the **kubectl logs amster-xxxxxxxx-yyyyy -c amster -f** command to stream the Amster pod's log to standard output.

The deployment clones the Git repository containing the initial AM configuration. Before the AM and DS servers become available, the following output appears:

```
. . .
+ ./amster-install.sh
Waiting for AM server at http://openam:80/openam/config/options.htm
Got Response code 000
response code 000. Will continue to wait
Got Response code 000
response code 000. Will continue to wait
Got Response code 000
. . .
```

When Amster starts to configure AM, the following output appears:

```
. . .
Got Response code 200
AM web app is up and ready to be configured
About to begin configuration
Executing Amster to configure AM
Executing Amster script /opt/amster/scripts/00_install.amster
Jul 31, 2019 4:51:29 PM java.util.prefs.FileSystemPreferences$1 run
INFO: Created user preferences directory.
Amster OpenAM Shell (6.5.2 build 314d553429, JVM: 1.8.0_151)
Type ':help' or ':h' for help
.
-----
am> :load /opt/amster/scripts/00_install.amster
07/31/2019 04:51:32:699 PM GMT: Checking license acceptance...
07/31/2019 04:51:32:700 PM GMT: License terms accepted.
07/31/2019 04:51:32:706 PM GMT: Checking configuration directory /home/forgerock/openam.
07/31/2019 04:51:32:707 PM GMT: ...Success.
07/31/2019 04:51:32:712 PM GMT: Tag swapping schema files.
07/31/2019 04:51:32:747 PM GMT: ...Success.
07/31/2019 04:51:32:747 PM GMT: Loading Schema odsee_config_schema.ldif
07/31/2019 04:51:32:825 PM GMT: ...Success.
07/31/2019 04:51:32:825 PM GMT: Loading Schema odsee_config_index.ldif
07/31/2019 04:51:32:855 PM GMT: ...Success.
07/31/2019 04:51:32:855 PM GMT: Loading Schema cts-container.ldif
07/31/2019 04:51:32:945 PM GMT: ...Success
.
. . .
```

The following output indicates that deployment is complete:

```

07/31/2019 04:51:53:215 PM GMT: Setting up monitoring authentication file.
Configuration complete!
Executing Amster script /opt/amster/scripts/01_import.amster
Amster OpenAM Shell (6.5.2 build 314d553429, JVM: 1.8.0_151)
Type ':help' or ':h' for help
.
-----
am> :load /opt/amster/scripts/01_import.amster
Importing directory /git/config/default/am/empty-import
Import completed successfully
Configuration script finished
+ pause
+ echo Args are 0
+ echo Container will now pause. You can use kubectl exec to run export.sh
+ true
+ sleep 1000000
Args are 0
Container will now pause. You can use kubectl exec to run export.sh

```

5. Delete the AM pod to force an AM server restart.

After the restart, AM uses the appropriate CTS configuration.

- a. Run the **kubectl get pods** command to get the AM pod's name.

The name of the AM pod starts with the string **prod-openam**.

- b. Delete the pod:

```
$ kubectl delete pod openam-pod
```

- c. Wait several seconds and then run the **kubectl get pods** command again. Observe that a new AM pod has started up.

When you installed the **openam** chart in Step 1, the AM server started up before the configuration was available. The **amster** chart then installed AM configuration, including the CTS configuration. But the CTS configuration is not hot-swappable; AM cannot change the CTS configuration without a restart.

## To Deploy IDM

1. Install the **postgres-openidm** chart:

```

$ helm install --name prod-postgres-openidm --namespace prod \
  --values common.yaml --values postgres-openidm.yaml /path/to/forgeops/helm/postgres-openidm
NAME: prod-postgres-openidm
LAST DEPLOYED: Thu Aug 1 12:41:53 2019
NAMESPACE: prod
STATUS: DEPLOYED

RESOURCES:
==> v1/PersistentVolumeClaim
NAME AGE

```

```

postgres-openidm 1s

==> v1/Service
postgresql 1s

==> v1beta1/Deployment
postgres-openidm 1s

==> v1/Pod(related)

NAME                                READY  STATUS   RESTARTS  AGE
postgres-openidm-6b6fb898cb-mthrj  0/1    Pending  0          1s

==> v1/Secret

NAME            AGE
postgres-openidm 1s

==> v1/ConfigMap
openidm-sql 1s

NOTES:
PostgreSQL can be accessed via port 5432 on the following DNS name from within your cluster:
prod-postgres-openidm-postgres-openidm.prod.svc.cluster.local

To get your user password run:

PGPASSWORD=$(printf $(printf '\%o' `kubectl get secret --namespace prod prod-postgres-openidm-postgres-openidm -o jsonpath="{.data.postgres-password[*]}"`);echo)

To connect to your database run the following command (using the env variable from above):

kubectl run prod-postgres-openidm-postgres-openidm-client --rm --tty -i --image postgres \
--env "PGPASSWORD=$PGPASSWORD" \
--command -- psql -U openidm \
-h prod-postgres-openidm-postgres-openidm postgres

```

## 2. Install the `openidm` chart:

```

$ helm install --name prod-openidm --namespace prod \
  --values common.yaml --values openidm.yaml /path/to/forgeops/helm/openidm
NAME:      prod-openidm
LAST DEPLOYED: Thu Aug  1 12:46:24 2019
NAMESPACE: prod
STATUS:    DEPLOYED

RESOURCES:
==> v1/Secret
NAME            AGE
openidm-secrets-env 2s
openidm-secrets  2s

==> v1/ConfigMap
prod-openidm-openidm 1s
idm-boot-properties  1s
idm-logging-properties 1s

==> v1/Service

```

```

openidm ls
==> v1beta1/StatefulSet
prod-openidm-openidm ls

==> v1beta1/Ingress
openidm ls

==> v1/Pod(related)

NAME                                READY  STATUS   RESTARTS  AGE
prod-openidm-openidm-0             0/2    Init:0/1  0          1s

NOTES:
OpenIDM should be available soon at the ingress address of http://openidm.prod.example.com

It can take a few minutes for the ingress to become ready.

```

### 3. Check the status of the pods in the deployment until all pods are ready:

#### a. Run the **kubect** **get pods** command:

```

$ kubectl get pods
NAME                                READY   STATUS    RESTARTS  AGE
amster-7b8b6d9667-hxq57            2/2     Running   0          31m
configstore-0                      1/1     Running   0          40m
configstore-1                      1/1     Running   0          40m
ctsstore-0                         1/1     Running   0          40m
ctsstore-1                         1/1     Running   0          39m
dsadmin-69fb874f74-dln6b           1/1     Running   0          41m
postgres-openidm-6b6fb898cb-mthrj  1/1     Running   0          17m
prod-openam-openam-6f846bbf69-867rc 1/1     Running   0          24m
prod-openidm-openidm-696685d8cb-pt5x4 2/2     Running   0          12m
userstore-0                        1/1     Running   0          40m
userstore-1                        1/1     Running   0          39m

```

#### b. Review the output. Deployment is complete when:

- The **READY** column indicates all containers are available. The entry in the **READY** column represents [total number of containers/number of available containers].
- All entries in the **STATUS** column indicate **Running**.

#### c. If necessary, continue to query your deployment's status until all the pods are ready.

## To Deploy IG

### 1. Install the **openig** chart:

```
$ helm install --name prod-openig --namespace prod \
--values common.yaml --values openig.yaml /path/to/forgeops/helm/openig
NAME:      prod-openig
LAST DEPLOYED: Thu Aug  1 12:46:24 2019
NAMESPACE: prod
STATUS: DEPLOYED

RESOURCES:
==> v1/Service
openig 1s

==> v1beta1/Deployment
prod-openig-openig 1s

==> v1beta1/Ingress
openig 1s

==> v1/Pod(related)
NAME                                READY  STATUS             RESTARTS  AGE
prod-openig-openig-9b4dfb574-jdp4n  0/1    ContainerCreating  0         0s

NOTES:
1. Get the application URL by running these commands:
  export POD_NAME=$(kubectl get pods --namespace prod -l "app=prod-openig-openig" -o jsonpath="{.items[0].metadata.name}")
  echo "Visit http://127.0.0.1:8080 to use your application"
  kubectl port-forward $POD_NAME 8080:8080

If you have an ingress controller, you can also access IG at
http://openig.prod.example.com
```

## 2. Check the status of the pods in the deployment until all pods are ready:

### a. Run the **kubectl get pods** command:

```
$ kubectl get pods
NAME                                READY  STATUS             RESTARTS  AGE
amster-7b8b6d9667-hxq57            2/2    Running            0         31m
configstore-0                      1/1    Running            0         40m
configstore-1                      1/1    Running            0         40m
ctsstore-0                         1/1    Running            0         40m
ctsstore-1                         1/1    Running            0         39m
dsadmin-69fb874f74-dln6b           1/1    Running            0         41m
postgres-openidm-6b6fb898cb-mthrj  1/1    Running            0         17m
prod-openam-openam-6f846bbf69-867rc 1/1    Running            0         24m
prod-openidm-openidm-696685d8cb-pt5x4 2/2    Running            0         12m
prod-openig-openig-9b4dfb574-jdp4n  1/1    Running            0         1m
userstore-0                       1/1    Running            0         40m
userstore-1                       1/1    Running            0         39m
```

### b. Review the output. Deployment is complete when:

- The **READY** column indicates all containers are available. The entry in the **READY** column represents [total number of containers/number of available containers].

- All entries in the **STATUS** column indicate **Running**.

- If necessary, continue to query your deployment's status until all the pods are ready.

## To Deploy the web Application

The **web** application is required for running the IG reverse proxy benchmark.

1. Install the **web** chart:

```
$ helm install --name prod-web --namespace prod \
--values common.yaml /path/to/forgeops/helm/web
Release "prod-web" does not exist. Installing it now.
NAME: prod-web
LAST DEPLOYED: Tue Jul 30 16:33:45 2019
NAMESPACE: prod
STATUS: DEPLOYED

RESOURCES:
==> v1/Pod(related)
NAME                                READY  STATUS   RESTARTS  AGE
prod-web-747bdd74c7-zz8wd           0/1    Pending  0          0s

==> v1/ConfigMap
NAME      AGE
prod-web  0s

==> v1/Service
prod-web  0s

==> v1beta2/Deployment
prod-web  0s

==> v1beta1/Ingress
prod-web  0s

NOTES:
1. Get the application URL by running these commands:
```

2. Check the status of the pods in the deployment until all pods are ready:

- Run the **kubectl get pods** command:



```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
amster-7b8b6d9667-hxq57	2/2	Running	0	31m
configstore-0	1/1	Running	0	40m
configstore-1	1/1	Running	0	40m
ctsstore-0	1/1	Running	0	40m
ctsstore-1	1/1	Running	0	39m
dsadmin-69fb874f74-dln6b	1/1	Running	0	41m
postgres-openidm-6b6fb898cb-mthrj	1/1	Running	0	17m
prod-openam-openam-6f846bbf69-867rc	1/1	Running	0	24m
prod-openidm-openidm-696685d8cb-pt5x4	2/2	Running	0	12m
prod-openig-openig-9b4dfb574-jdp4n	1/1	Running	0	1m
prod-web-747bdd74c7-zz8wd	1/1	Running	0	5s
userstore-0	1/1	Running	0	40m
userstore-1	1/1	Running	0	39m

- b. Review the output. Deployment is complete when:
  - The **READY** column indicates all containers are available. The entry in the **READY** column represents [total number of containers/number of available containers].
  - All entries in the **STATUS** column indicate **Running**.
- c. If necessary, continue to query your deployment's status until all the pods are ready.

## To Scale the Deployment

The ForgeRock Identity Platform Helm charts create one AM pod, one IDM pod, and two IG pods. To create a second pod for AM and IDM, run the **kubectl scale replicas=2** command.

1. Get the AM deployment name:

```
$ kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
amster	1	1	1	1	34m
dsadmin	1	1	1	1	45m
postgres-openidm	1	1	1	1	20m
prod-openam-openam	1	1	1	1	35m
prod-openidm-openig	1	1	1	1	16m

2. Scale the number of AM pods to two:

```
$ kubectl scale --replicas=2 deployment prod-openam-openam
deployment.extensions/prod-openam-openam scaled
```

3. Get the IDM stateful set name:

```
$ kubectl get statefulsets
```

NAME	DESIRED	CURRENT	AGE
prod-openam-openidm	1	1	35m
userstore	1	1	54m
configstore	1	1	55m
ctsstore	1	1	56m

4. Scale the number of IDM pods to two:

```
$ kubectl scale --replicas=2 statefulset prod-openidm-openidm
statefulset.extensions/prod-openidm-openidm scaled
```

5. Verify that two **openam**, **openidm**, and **openig** pods are available in your cluster:

```
$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
amster-7b8b6d9667-hxq57	2/2	Running	0	37m
configstore-0	1/1	Running	0	47m
configstore-1	1/1	Running	0	46m
ctsstore-0	1/1	Running	0	46m
ctsstore-1	1/1	Running	0	46m
dsadmin-69fb874f74-dln6b	1/1	Running	0	48m
postgres-openidm-6b6fb898cb-mthrx	1/1	Running	0	23m
prod-openam-openam-6f846bbf69-867rc	1/1	Running	0	30m
prod-openam-openam-6f846bbf69-bbpfq	1/1	Running	0	1m
prod-openidm-openidm-696685d8cb-nbd2n	2/2	Running	0	27s
prod-openidm-openidm-696685d8cb-pt5x4	2/2	Running	0	19m
prod-openig-openig-6a12fe4501-n1402	2/2	Running	0	50s
prod-openig-openig-6a12fe4501-g2x01	2/2	Running	0	48m
userstore-0	1/1	Running	0	46m
userstore-1	1/1	Running	0	46m

## Chapter 4

# Using the CDM

This chapter shows you how to access and monitor the ForgeRock Identity Platform components that make up the CDM.

## 4.1. Accessing ForgeRock Identity Platform Services

This section shows you how to access ForgeRock Identity Platform components within the CDM.

AM, IDM, and IG are configured for access through the CDM cluster's Kubernetes ingress controller. You can access these components using their normal interfaces:

- For AM, the console and REST APIs.
- For IDM, the Admin UI and REST APIs.
- For IG, HTTP URLs. Note that because the CDM deploys IG in production mode, IG Studio is not available.

DS cannot be accessed through the ingress controller, but you can use Kubernetes methods to access the DS pods.

For more information about how AM, IDM, and IG have been configured in the CDM, see [the 6.5 CDM README file in the `forgeops-init` repository](#).

### 4.1.1. Before You Begin

The URLs you use to access ForgeRock Identity Platform services must be resolvable from your local computer.

If DNS does not resolve the following hostnames, add entries for them to your `/etc/hosts` file:

- `login.prod.my-domain`
- `openidm.prod.my-domain`
- `openig.prod.my-domain`

### 4.1.2. Accessing AM Services

Access the AM console and REST APIs as follows:

- "To Access the AM Console"
- "To Access the AM REST APIs"

### *To Access the AM Console*

1. By default, the `amster` Helm chart dynamically generates a random password for the `amadmin` user. It stores the password in the `amster-config` configmap.

To obtain the password, look for the `adminPwd` value in the `amster-config` configmap:

```
$ kubectl get configmaps amster-config -o yaml | grep adminPwd
--adminPwd "74xW6IShJF" \
```

2. (Optional) Delete the `amster-config` configmap so that the `amadmin` password is not publicly available.

```
$ kubectl delete configmaps amster-config
```

3. Open a new window or tab in a web browser.
4. Navigate to the AM deployment URL, `https://login.prod.my-domain/XUI/?service=adminconsole`.

The Kubernetes ingress controller handles the request, routing it to a running AM instance.

AM prompts you to log in.

5. Log in to AM as the `amadmin` user.

### *To Access the AM REST APIs*

1. Start a terminal window session.
2. If you have not already done so, obtain the password for the `amadmin` user. To obtain the password, perform Step 1 in "To Access the AM Console".
3. Run a `curl` command to verify that you can access the REST APIs through the ingress controller. For example:

```
$ curl \
  --request POST \
  --insecure \
  --header "Content-Type: application/json" \
  --header "X-OpenAM-Username: amadmin" \
  --header "X-OpenAM-Password: amadmin password" \
  --header "Accept-API-Version: resource=2.0" \
  --data "{}" \
  'https://login.prod.my-domain/json/realms/root/authenticate?
service=ldapService&authIndexType=service&authIndexValue=ldapService'
{
  "tokenId": "AQIC5wM2...",
  "successUrl": "/console",
  "realm": "/"
}
```

### 4.1.3. Accessing IDM Services

In the CDM, IDM is deployed with a read-only configuration. You can access the IDM Admin UI and REST APIs, but you cannot use them to modify the IDM configuration.

Access the IDM Admin UI and REST APIs as follows:

- "To Access the IDM Admin UI Console"
- "To Access the IDM REST APIs"

#### *To Access the IDM Admin UI Console*

1. Open a new window or tab in a web browser.
2. Navigate to the IDM Admin UI URL, `https://openidm.prod.my-domain/admin`.

The Kubernetes ingress controller handles the request, routing it to a running IDM instance.

IDM prompts you to log in.

3. Log in to the IDM Admin UI as the `openidm-admin` user with password `openidm-admin`.

#### *To Access the IDM REST APIs*

1. Start a terminal window session.
2. Run a **curl** command to verify that you can access the REST APIs through the ingress controller. For example:

```
$ curl \
  --request GET \
  --insecure \
  --header "X-OpenIDM-Username: openidm-admin" \
  --header "X-OpenIDM-Password: openidm-admin" \
  --data "{}" \
  https://openidm.prod.my-domain/openidm/info/ping
{
  "_id": " ",
  "_rev": "",
  "shortDesc": "OpenIDM ready",
  "state": "ACTIVE_READY"
}
```

#### 4.1.4. Accessing IG Services

In the CDM, IG is deployed in production mode, and therefore, IG Studio is disabled. You can access IG using a browser to verify that it is operational:

##### *To Verify IG Is Operational*

1. Open a new window or tab in a web browser.
2. Navigate to `https://openig.prod.my-domain`.

The Kubernetes ingress controller handles the request, routing it to IG.

You should see a message similar to the following:

```
hello and Welcome to OpenIG. Your path is /. OpenIG is using the default handler for this route.
```

#### 4.1.5. Accessing DS

The DS pods in the CDM are not exposed outside of the cluster. If you need to access one of the DS pods, use a standard Kubernetes method:

- Execute shell commands in DS pods using the **kubecttl exec** command.
- Forward a DS pod's LDAP port (1389) to your local computer. Then you can run LDAP CLI commands, for example **ldapsearch**. You can also use an LDAP editor such as Apache Directory Studio to access the directory.

For all directory pods in the CDM, the directory manager entry is `cn=Directory Manager` and the password is `password`.

## 4.2. Monitoring the CDM

This section contains procedures for accessing Grafana dashboards and the Prometheus web UI:

- "To Access Grafana Dashboards"
- "To Access the Prometheus Web UI"

### To Access Grafana Dashboards

1. Forward port 3000 on your local computer to port 3000 on the Grafana web server:

```
$ kubectl \
port-forward \
$(kubectl get pods --selector=app=grafana \
--output=jsonpath="{.items..metadata.name}" --namespace=monitoring) \
3000 --namespace=monitoring
Forwarding from 127.0.0.1:3000 -> 3000
Forwarding from [::1]:3000 -> 3000
```

2. In a web browser, navigate to <http://localhost:3000> to start the Grafana user interface.  
For information about the Grafana UI, see the [Grafana documentation](#).
3. When you're done using the Grafana UI, enter Cntl+c in the terminal window to stop port forwarding.

### To Access the Prometheus Web UI

1. Forward port 9090 on your local computer to port 9090 on the Prometheus web server:

```
$ kubectl \
port-forward \
$(kubectl get pods --selector=app=prometheus \
--output=jsonpath="{.items..metadata.name}" --namespace=monitoring) \
9090 --namespace=monitoring
Forwarding from 127.0.0.1:9090 -> 9090
Forwarding from [::1]:9090 -> 9090
```

2. In a web browser, navigate to <http://localhost:9090>.  
The Prometheus web UI appears in the browser.  
To use the Prometheus web UI, see the [Prometheus documentation](#).
3. When you're done using the Prometheus web UI, enter Cntl+c in the terminal window to stop port forwarding.

To enable monitoring in the CDM, see:

- "Deploying Monitoring Infrastructure" in the *Cloud Deployment Model Cookbook for GKE*
- "Deploying Monitoring Infrastructure"

For a description of the CDM monitoring architecture and information about how to customize CDM monitoring, see:

- *"Monitoring Your Deployment" in the Site Reliability Guide for GKE*
- *"Monitoring Your Deployment" in the Site Reliability Guide for Amazon EKS*



## Chapter 5

# Benchmarking the CDM Performance

This chapter provides steps you can take to replicate the Cloud Deployment Team benchmarking process. Topics include the following:

- "About CDM Benchmarking"
- "Before You Begin"
- "Running Directory Services Benchmark Tests"
- "Running AM, IDM, or IG Benchmark Tests"

## 5.1. About CDM Benchmarking

We've published benchmarking instructions and results to give you a means for validating your own CDM deployment. If you follow the benchmarking instructions *with no configuration modifications*, you should attain similar or better throughput and latency results.

### Important

These results are based on Cloud Deployment Model (CDM) benchmark tests we conducted through January, 2019. We conducted the tests using Kubernetes version 1.11, and default sizing and configurations described in this guide. Factors beyond the scope of the CDM or its default sizing and configuration may affect your benchmark test results. These factors might include (but are not limited to): updates to the cloud provider's SDK; changes to third-party software required for Kubernetes; changes you have made to sizing or configuration to suit your business needs.

### 5.1.1. Default Sizing and Configuration

By default, the CDM is sized to strike a balance between optimal performance and low production cost. After you've successfully deployed the CDM, your deployment is sized and configured to the following specifications.

Description	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Kubernetes version	1.11.2	1.11.2	1.11.2
Nodes	4	4	4

Description	Small Cluster	Medium Cluster	Large Cluster
vCPU	4	16	32
Memory	16 GB	64 GB	72 GB
Storage	gp2 100 GB	gp2 500 GB and 850 GB	gp2 500 GB and 1TB
Virtual Machine Size	m5.xlarge	m5.4xlarge	c5.9xlarge
Average DS user entry size	1K	1K	1K
AM CTS storage scheme	GrantSet	GrantSet	GrantSet
AM CTS reaper	Disabled on AM Enabled TTL on CTS DS	Disabled on AM Enabled TTL on CTS DS	Disabled on AM Enabled TTL on CTS DS

### 5.1.2. Cost Estimates for Running CDM on Amazon Web Services

The following table compares approximate monthly costs of running the CDM using the default small and medium worker node configurations. These approximations were derived using the AWS simple monthly calculator.

	Small Cluster	Medium Cluster	Large Cluster
<b>Number of Users</b>	1,000,000	10,000,000	100,000,000
<b>Benchmark Tests Date</b>	January, 2019	January, 2019	January, 2019
<b>Virtual Machines</b> Cost per VM. Reserved instances with standard 1-year term.	m5.xlarge (4x16) 4 nodes = \$359/month	m5.4xlarge (16x64) 4 nodes = \$1433/month	c5.9xlarge (32x72) 4 nodes = \$2814/month
<b>Network SSD Storage EBS</b> For databases. gp2 \$0.10 per GB Io1 \$0.125 per GB and \$0.065 per IOPS	400 GB gp2 400*0.1=\$40/month	3000 GB gp2 3000*0.1=\$300/month	4000 GB gp2 4000 GB = \$825/month
<b>Amazon EFS Storage</b> For backup. Bursting Throughput \$0.30/GB	\$300/month	\$600/month	\$1200/month
<b>Amazon S3 Cloud Storage</b> For backup archival. \$0.023/GB	1 TB = \$23/month	2 TB = \$46/month	4 TB = \$92/month
<b>Networking</b> Includes load balancing, zone data transfer, elastic IP, and other resources. Based on network traffic volume.	\$20/month	\$40/month	\$80/month
<b>Cloud SQL - RDS</b>	db.t2.medium	db.m5.large	db.m5.xlarge

	Small Cluster	Medium Cluster	Large Cluster
This is optional if IDM is deployed. Multi-Availability zone and Multi-AZ storage.	100GB storage 200GB backup 105+40=\$145/month	200GB storage 400GB backup 256+80=\$336/month	300GB storage 600 GB backup 512+160=\$672/month
<b>Cloud Logging and Monitoring</b> Optional CloudWatch dashboard.	\$100/month	\$200/month	\$300/month
<b>Kubernetes Control Plane Usage</b>	\$144/month	\$144/month	\$144/month
<b>Estimated Total Cost Per Month</b>	\$563 \$1135 with all options.	\$1917 \$3099 with all options.	\$3863 \$6127 with all options.

### 5.1.3. How CDM Benchmarking Relates to Your Deployment

After you've successfully deployed the CDM, run the tests in this chapter. Once you've validated your own results against the benchmarks reported here, you can be sure that ForgeRock Identity Platform is successfully running in the cloud to CDM specifications.

The CDM is designed to:

- Conform to DevOps best practices.
- Facilitate continuous integration and continuous deployment.
- Scale and deploy on any Kubernetes environment in the cloud.

If you require higher performance than the benchmarks reported here, you can scale your deployment horizontally and vertically. Vertically scaling ForgeRock Identity Platform works particularly well in the cloud. For more information about scaling your deployment, contact your qualified ForgeRock partner or technical consultant.

## 5.2. Before You Begin

Before you can test CDM performance, you must generate test users, provision the CDM userstores, and prime the Directory Servers. Provision the first userstore by creating and importing an LDIF file. Provision the second userstore using the DS backup and restore capabilities and setting the generation ID. This ensures that the data is synchronized for subsequent replication.

Complete the following steps:

### *To Provision the Userstores with Test Users*

1. Open a shell in the `userstore-0` pod. Example:

```
$ kubectl exec userstore-0 -it bash
```

2. Edit the `/opt/openssh/config/MakeLDIF/example.template` file:
  - a. Change the `suffix` definition on the first line to `define suffix=ou=identities`.
  - b. Change the `numusers` definition:
    - For a small cluster, specify `define numusers=1000000`.
    - For a medium cluster, specify `define numusers=10000000`.
    - For a large cluster, specify `define numusers=100000000`.
  - c. In the `branch: [suffix]` section, change the line `objectClass: domain` to `objectClass: organizationalUnit`.

An example of the edited file snippet for a small cluster:

```
define suffix=ou=identities
define maildomain=example.com
define numusers=1000000

branch: [suffix]
objectClass: top
objectClass: organizationalUnit
...
```

3. Provision the `userstore-0` directory with test users:
  - a. Generate test users using the `makeldif` command:
 

```
$ /opt/openssh/bin/makeldif -o /tmp/lm.ldif \
/opt/openssh/config/MakeLDIF/example.template
Processed 1000 entries
Processed 2000 entries
Processed 3000 entries
Processed 4000 entries
...
LDIF processing complete. 1000002 entries written
```
  - b. Zip the resulting LDIF file, and copy the zipped archive to the `/opt/openssh/bak` directory:
 

```
$ gzip /tmp/lm.ldif
$ cp /tmp/lm.ldif.gz /opt/openssh/bak
```
  - c. Import the test users into the `userstore-0` directory:

```
$ /opt/openssl/bin/import-ldif -F -c -l /opt/openssl/bak/1m.ldif.gz \
-p 4444 -D "cn=Directory Manager" -w password --trustAll -n amIdentityStore
Import task 20190731145319759 scheduled to start immediately
[31/Jul/2019:14:53:19 +0000] severity="NOTICE" msgCount=0 msgID=org.opens
.messages.backend-413 message="Import task 20190731145319759 started
execution"
...
[31/Jul/2019:14:53:43 +0000] severity="NOTICE" msgCount=14 msgID=org.opens.messages.backend-531
message="Total import time was 22 seconds. Phase one processing completed in 14 seconds, phase
two processing completed in 8 seconds"
[31/Jul/2019:14:53:43 +0000] severity="NOTICE" msgCount=15 msgID=org.opens
.messages.backend-519 message="Processed 100002 entries, imported 100002,
skipped 0, rejected 0 and migrated 0 in 23 seconds (average rate 4288.4/
sec)"
...
Import task 20190731145319759 has been successfully completed
```

4. For small, medium, and large clusters, prime the `userstore-0` directory. Run the following LDAP search:

```
$ ldapsearch -p 1389 -D "cn=Directory Manager" \
-w password -b "ou=identities" objectclass=* > /dev/null
```

It could take as long as ten minutes or more to complete the search operation.

5. Make a backup of the `userstore-0` directory containing the imported test users:

```
$ /opt/openssl/scripts/backup.sh
Backup task 20190731145523468 scheduled to start immediately
[31/Jul/2019:14:55:23 +0000] severity="NOTICE" msgCount=0 msgID=org.opens.messages.backend-413
message="Backup task 20190731145523468 started execution"
[31/Jul/2019:14:55:23 +0000] severity="NOTICE" msgCount=1 msgID=org.opens.messages.tool-280
message="Starting backup for backend amIdentityStore"
[31/Jul/2019:14:55:30 +0000] severity="NOTICE" msgCount=2 msgID=org.opens.messages.utility-321
message="Archived backup file: 00000000.jdb"
[31/Jul/2019:14:55:30 +0000] severity="NOTICE" msgCount=3 msgID=org.opens.messages.tool-283
message="The backup process completed successfully"
[31/Jul/2019:14:55:30 +0000] severity="NOTICE" msgCount=4 msgID=org.opens.messages.backend-414
message="Backup task 20190731145523468 finished execution in the state Completed successfully"
Backup task 20190731145523468 has been successfully completed
```

6. Close the shell in the `userstore=0` pod.

```
$ exit
```

7. Provision the `userstore-1` directory with test users:

- a. Open a shell in the `userstore-1` pod:

```
$ kubectl exec userstore-1 -it bash
```

- b. Restore the test users into the `userstore-1` directory from the `userstore-0` directory backup:

```
$ /opt/openssl/bin/restore \
--hostname localhost --port 4444 --bindDN cn="Directory Manager" \
--bindPasswordFile /var/run/secrets/openssl/dirmanager.pw --trustAll \
--backupDirectory /opt/openssl/bak/default/userstore-prod/amIdentityStore
Restore task 20190731150155326 scheduled to start immediately
[31/Jul/2019:15:01:55 +0000] severity="NOTICE" msgCount=0 msgID=org.openss.messages.backend-413
message="Restore task 20190731150155326 started execution"
[31/Jul/2019:15:01:55 +0000] severity="NOTICE" msgCount=1 msgID=org.openss.messages.backend-370
message="The backend amIdentityStore is now taken offline"
[31/Jul/2019:15:01:56 +0000] severity="NOTICE" msgCount=2 msgID=org.openss.messages.utility-320
message="Restored backup file: 00000000.jdb (size 139284948)"
[31/Jul/2019:15:01:56 +0000] severity="INFORMATION" msgCount=3 msgID=org.openss.messages.backend
-410 message="JE backend 'amIdentityStore' does not specify the number of cleaner threads:
defaulting to 2 threads"
[31/Jul/2019:15:01:56 +0000] severity="INFORMATION" msgCount=4 msgID=org.openss.messages.backend
-411 message="JE backend 'amIdentityStore' does not specify the number of lock tables: defaulting
to 13"
[31/Jul/2019:15:01:58 +0000] severity="NOTICE" msgCount=5 msgID=org.openss.messages.backend-513
message="The database backend amIdentityStore containing 100002 entries has started"
[31/Jul/2019:15:01:58 +0000] severity="WARNING" msgCount=6 msgID=org.openss.messages.replication
-96 message="Directory server DS(11) has connected to replication server RS(11) for domain
"ou=identities" at 10.32.3.11:8989, but the generation IDs do not match, indicating that a full
re-initialization is required. The local (DS) generation ID is 19094517 and the remote (RS)
generation ID is 65525"
[31/Jul/2019:15:01:58 +0000] severity="NOTICE" msgCount=7 msgID=org.openss.messages.backend-414
message="Restore task 20190731150155326 finished execution in the state Completed successfully"
Restore task 20190731150155326 has been successfully completed
```

8. Set the generation ID on the `userstore-1` server so that replication between the `userstore-0` and `userstore-1` servers works:

- a. Review the standard output from the `restore` command that you ran in the preceding step. Locate the message similar to this message:

```
...the generation IDs do not match, indicating that a full re-initialization is required. The
local (DS) generation ID is 19094517 and the remote (RS) generation ID is 65525"
```

- b. Obtain the local generation ID from the message. In the example, the local generation ID is `19094517`.
- c. Run the following command to set the local generation ID of the `userstore-1` server to the generation ID of the `userstore-1` server:

```
$ /opt/openssl/bin/ldapmodify -p 1389 -D "cn=Directory Manager" -w password <<EOF
dn: ds-task-id=mytask-1234,cn=Scheduled Tasks,cn=Tasks
objectClass: ds-task-reset-generation-id
ds-task-class-name: org.openss.server.tasks.SetGenerationIdTask
ds-task-reset-generation-id-domain-base-dn: ou=identities
ds-task-reset-generation-id-new-value: 19094517
EOF
```

In the preceding command:

- The value for `ds-task-id` can be any unique value.

- The value for `ds-task-reset-generation-id-new-value` must be the local generation ID that you obtained from the **restore** command output.

- Review the `/opt/openssl/logs/errors` file. You should see a message similar to the following:

```
[02/Aug/2019:18:08:54 +0100] category=SYNC severity=NOTICE msgID=78 msg=The generation ID for domain "dc=example,dc=com" has been reset to 65525
```

- Continue examining the `/opt/openssl/logs/errors` file. After a short wait, you should see a message similar to the following:

```
[31/Jul/2019:17:56:12 +0100] category=SYNC severity=NOTICE msgID=62 msg=Directory server DS(7775) has connected to replication server RS(20607) for domain "dc=example,dc=com" at 127.0.0.1:8990 with generation ID 65525
```

- Prime the `userstore-1` directory by running the following LDAP search:

```
$ ldapsearch -p 1389 -D "cn=Directory Manager" \
-w password -b "ou=identities" objectclass=* > /dev/null
```

It could take as long as ten minutes or more to complete the search operation.

After both of the DS servers have been primed, you can run any of the benchmark tests in this guide.

## 5.3. Running Directory Services Benchmark Tests

The Cloud Deployment Team conducted the DS tests using the `searchrate`, `modrate`, `addrate` binaries. These binaries are packaged with DS. See the `ds-bench.sh` script for details such as number of client threads, duration, and filters.

### To Run DS Benchmark Tests

- The userstore must be provisioned, and the Directory Server must be primed.

See "To Provision the Userstores with Test Users".

- Open a shell in the `userstore-1` pod, and run the `ds-bench.sh` script. Example:

```
$ scripts/ds-bench.sh srch 600 localhost 1000000
```

In this example:

- `srch` specifies the `searchrate` tool.
- `600` specifies the duration (in seconds) of the load.
- `localhost` specifies the target host.

- 1000000 specifies the number of users.

See the `ds-bench` script in Git at [forgeops/docker/ds/scripts/ds-bench.sh](https://github.com/forgeops/docker/ds/scripts/ds-bench.sh).

### 5.3.1. Searchrate

The DS `searchrate` utility measures the search throughput and response time of a directory service using user-defined searches. Before running this test, the userstore must be provisioned with users. See "Before You Begin" for more information.

To run the `searchrate` utility, open a shell in the `userstore-1` pod and run the `ds-bench` script. Example:

```
$ /opt/openshift/scripts/ds-bench.sh srch 600 localhost 1000000
```

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Throughput	~17 k searches/sec	~46 k searches/sec	~24 k searches/sec
Response times	Mean: 31 ms	Mean:11 ms	Mean:11 ms
Resource requests	Memory: 3 Gi CPU: 2000 m	Memory: 8 Gi CPU: 2000 m	Memory: 20 Gi CPU: 4000 m
Resource limits	Memory: 5 Gi CPU: 3000 m	Memory: 11 Gi CPU: 8000 m	Memory: 27 Gi CPU: 10000 m
Pod usage	Memory: 4590 Mi CPU: 2000 m	Memory: 9868 Mi CPU: 7946 m	Memory: 25299 Mi CPU: 7177 m
Node usage	Memory: 13400 Mi CPU: 2070 m	Memory: 20264 Mi CPU: 8017 m	Memory: 47709 Mi CPU: 12384 m

### 5.3.2. Modrate

The DS `modrate` utility measures the modify throughput and response time of a directory service using user-defined modifications. Before running this test, the userstore must be provisioned with users. See "Before You Begin" for more information.

To run the `modrate` utility, open a shell in the `userstore-1` pod and run the `ds-bench` script. Example:

```
$ /opt/openshift/scripts/ds-bench.sh mod 600 localhost 1000000
```

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Throughput	~2.4 K mods/sec	~3.3 K mods/sec	~5 K mods/sec



Metric	Small Cluster	Medium Cluster	Large Cluster
Response times	Mean: 5 ms	Mean: 3 ms	Mean: 2 ms
Resource requests	Memory: 3 Gi CPU: 2000 m	Memory: 8 Gi CPU: 2000 m	Memory: 20 Gi CPU: 4000 m
Resource limits	Memory: 5 Gi CPU: 3000 m	Memory: 11 Gi CPU: 8000 m	Memory: 27 Gi CPU: 10000 m
Pod usage	Memory: 5000 Mi CPU: 1980 m	Memory: 9723 Mi CPU: 9428 m	Memory: 27276 Mi CPU: 9893 m
Node usage	Memory: 13800 Mi CPU: 2060 m	Memory: 20102 Mi CPU: 12004 m	Memory: 47709 Mi CPU: 12384 m

### 5.3.3. Addrate

The DS **addrate** utility measures the add throughput and response time of a directory server using user-defined entries. Before running this test, the userstore must be provisioned with users. See "Before You Begin" for more information.

To run the **addrate** utility, open a shell in the **userstore-1** pod and run the **ds-bench** script. Example:

```
$ /opt/opensj/scripts/ds-bench.sh add 600 localhost 1000000
```

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Throughput	~300 adds/sec	~600 adds/sec	~214 adds/sec
Response times	Mean: 12 ms	Mean: 9 ms	Mean: 37 ms
Resource requests	Memory: 3 Gi CPU: 2000 m	Memory: 8 Gi CPU: 2000 m	Memory: 20 Gi CPU: 4000 m
Resource limits	Memory: 5 Gi CPU: 3000 m	Memory: 11 Gi CPU: 8000 m	Memory: 27 Gi CPU: 10000 m
Pod usage	Memory: 5000 Mi CPU: 1990 m	Memory: 10060 Mi CPU: 4291 m	Memory: 27112 Mi CPU: 4889 m
Node usage	Memory: 13900 Mi CPU: 2330 m	Memory: 20634 Mi CPU: 9514 m	Memory: 47709 Mi CPU: 12384 m

## 5.4. Running AM, IDM, or IG Benchmark Tests

The Cloud Deployment Team developed the **gatling-benchmark** Helm chart for benchmarking AM, IDM, and IG performance. First the Helm chart runs one of a number of Scala-based use case simulations. Then the Helm chart exposes test results that you can view in a browser or download in a zipped file.

## To Run AM, IDM, or IG Benchmark Tests

1. In the `/path/to/forgeops/helm/gatling-benchmark` directory, edit the `values.yaml` file.
  - a. In the `testname` field, specify the name of the simulation you want to run.  
See "Access Management Benchmark Tests", "Identity Management Benchmark Tests", or "Identity Gateway Benchmark Tests" for more information.
  - b. In the `users` field, specify `1000000` for a small cluster, or `10000000` for a medium cluster. You can specify `100000000` for a large AM or a large IG cluster.
  - c. In the `concurrency` field, specify `50`.
  - d. In the `duration` field, specify `600`.
2. The userstores must be provisioned, and the Directory Server must be primed.  
See "To Provision the Userstores with Test Users".
3. To start the benchmark test, install the `gatling-benchmark` Helm chart. By default, the test begins immediately. Example:

```
$ helm install --name benchmark helm/gatling-benchmark
NAME: benchmark
LAST DEPLOYED: Tue Jul 30 14:56:40 2019
NAMESPACE: prod
STATUS: DEPLOYED

RESOURCES:
==> v1beta1/Ingress
NAME      AGE
gatling  1s

==> v1/Pod(related)

NAME                                READY  STATUS   RESTARTS  AGE
forgeops-benchmark-5957478bb4-qd9d5  0/1    Pending  0          0s

==> v1/ConfigMap

NAME      AGE
nginx-conf 1s

==> v1/PersistentVolumeClaim
forgeops-benchmark-pvc 1s

==> v1/Service
benchmark-gatling-benchmark 1s

==> v1beta1/Deployment
forgeops-benchmark 1s
```

- To verify that the **gatling-benchmark** Helm chart is installed, run the **kubectl get pods** command. For example:

```
$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
amster-689d795777-59p67             2/2     Running   0           1d
configstore-0                       1/1     Running   0           8h
configstore-1                       1/1     Running   0           9h
ctsstore-0                          1/1     Running   0           8h
ctsstore-1                          1/1     Running   0           1d
dsadmin-784fd647b4-pzbc5            1/1     Running   0           1d
forgeops-benchmark-5957478bb4-qd9d5 0/1     Init:0/2   0           3h
postgres-openidm-6fbf98d5bc-v2hbm   1/1     Running   0           1d
prod-openam-openam-5cf6b9498-c4sgh  1/1     Running   0           8h
prod-openam-openam-5cf6b9498-cp247  1/1     Running   1           9h
prod-openidm-openidm-0              2/2     Running   1           9h
prod-openidm-openidm-1              2/2     Running   1           1d
userstore-0                         1/1     Running   0           9h
userstore-1                         1/1     Running   0           8h]
```

In this example, the **forgeops-benchmark-5957478bb4-qd9d5** pod is created and initializing.

- In the **/path/to/forgeops/helm/gatling-benchmark** directory, run the **get-logs.sh** script.

Gatling output is sent directly to the console. For example:

```
$ ./get-logs.sh
[Loop n.0
  READY to run gatling.sh

-----
Provided script parameters:
java_options: -Dam_host=login.prod.perf.forgerock-qa.com -Dam_port=443 -Dam_protocol=https -
Didm_host=openidm.prod.perf.forgerock-qa.com -Didm_port=443 -Didm_protocol=https -Dig_host=openig
-Dig_port=80 -Dig_protocol=http -DlogoutPercent=0 -Dusers=1000000 -Dconcurrency=50 -Dduration=30 -
Dwarmup=60 -Dissue_token_info=false -Doauth2_client_id=client0IDC_0 -Doauth2_client_pw=password
cmd_options:
gatling_args:-m -s am.AMAccessTokenSim -rd am.AMAccessTokenSim
-----

Starting gatling simulation

-----
GATLING COMMAND: gatling.sh -m -s am.AMAccessTokenSim -rd am.AMAccessTokenSim
-----
GATLING_HOME is set to /opt/gatling
20:11:29.149 [WARN ] i.g.c.ZincCompiler$ - Pruning sources from previous analysis, due to
incompatible CompileSetup.
Simulation am.AMAccessTokenSim started...

=====
2018-11-27 20:11:42                                5s elapsed
---- Requests -----
> Global (OK=389 K0=0 )
> Rest login stage (OK=131 K0=0 )
> Authorize stage (OK=130 K0=0 )
> AccessToken stage (OK=128 K0=0 )
```

```

---- OAuth2 Auth code flow -----
[-----] 0%
waiting: 46 / active: 4 / done:0
=====

=====
2018-11-27 20:11:47 10s elapsed
---- Requests -----
> Global (OK=2178 K0=0 )
> Rest login stage (OK=727 K0=0 )
> Authorize stage (OK=726 K0=0 )
> AccessToken stage (OK=725 K0=0 )

---- OAuth2 Auth code flow -----
[-----] 0%
waiting: 42 / active: 8 / done:0
=====

=====
2018-11-27 20:11:52 15s elapsed
---- Requests -----
> Global (OK=5506 K0=0 )
> Rest login stage (OK=1840 K0=0 )
> Authorize stage (OK=1834 K0=0 )
> AccessToken stage (OK=1832 K0=0 )

---- OAuth2 Auth code flow -----
[-----] 0%
waiting: 38 / active: 12 / done:0
=====

=====
2018-11-27 20:11:57 20s elapsed
---- Requests -----
> Global (OK=10577 K0=0 )
> Rest login stage (OK=3532 K0=0 )
> Authorize stage (OK=3525 K0=0 )
> AccessToken stage (OK=3520 K0=0 )

---- OAuth2 Auth code flow -----
[-----] 0%
waiting: 34 / active: 16 / done:0
=====

```

The test will continue to run for the duration you specified in the `values.yaml` file.

6. (Optional) To stop a test, run the **helm del --purge** command. Example:

```
$ helm del --purge benchmark
```

## To View Test Results in a Browser


1. Confirm that an ingress address with an FQDN exists in the `/etc/hosts` file. Example:

```
$ kubectl get ingress
NAME          HOSTS                ADDRESS          PORTS    AGE
gatling       gatling.prod.forgeops.com 35.231.229.20    80       4h
openam        login.prod.forgeops.com   35.231.229.20    80, 443  1d
openidm       openidm.prod.forgeops.com 35.231.229.20    80, 443  1d
```

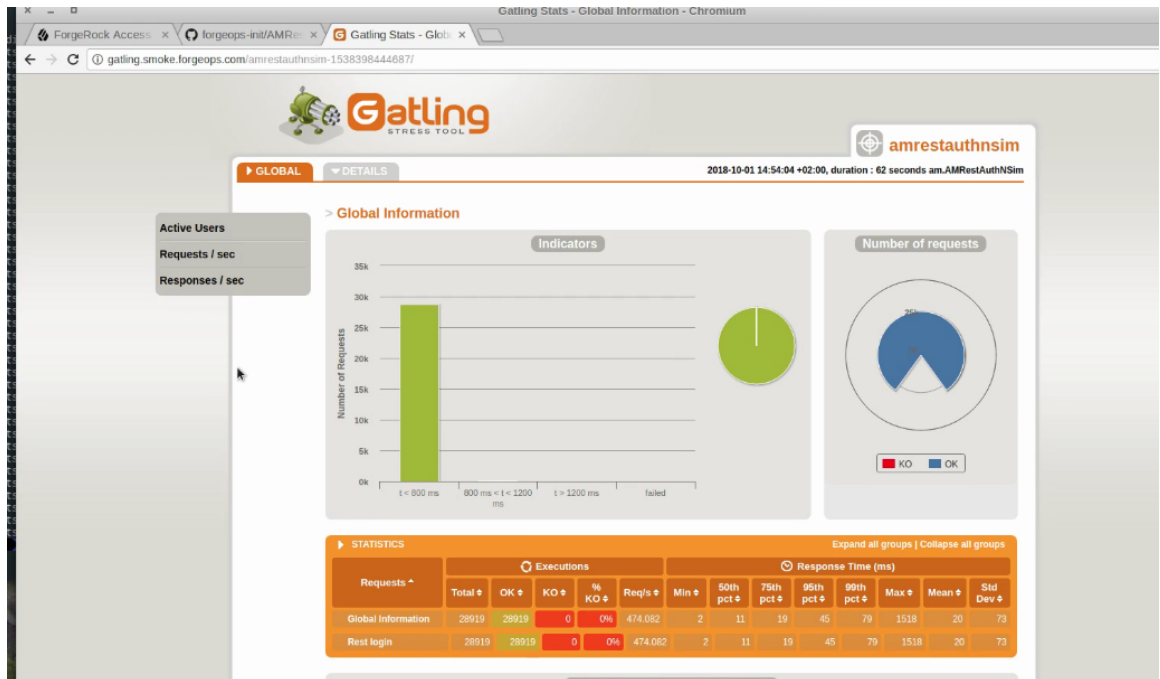
2. In a browser go to the URL of the Gatling host listed in the previous step.

In this example, the Gatling host URL is <http://gatling.prod.forgeops.com>.

- a. You can download the HTML report in the `*.tar.gz` file for future use. Once the `gatling-benchmark` Helm chart is deleted, you can no longer access these files.



- b. Click the link to go to the Gatling dashboard.



## 5.4.1. Access Management Benchmark Tests

The `AMRestAuthNSim.scala` simulation tests the authentication rates using the REST API. This simulation was run for stateless and stateful session tokens. A stateful session has an opaque token that persists in the CTS server so that the session token can be invalidated immediately if needed.

The `AMAccessTokenSim.scala` simulation tests stateless and stateful tokens in the OAuth 2.0 authorization code flow. A stateless token is a JWT that carries its state as a cookie set in the browser; the JWT has no corresponding server component.

### 5.4.1.1. REST Login - Stateful Authentication

This benchmark test measures throughput and response times of an AM server performing stateful authentications.

Before running the benchmark test:

1. Make sure the userstore is provisioned, and the Directory Services cache is primed.

See "To Provision the Userstores with Test Users".

2. Configure AM for stateful sessions:
  - a. Log in to the AM console as the `amadmin` user. For details, see "Accessing AM Services".
  - b. Select the top-level realm.
  - c. Select Properties.
  - d. Disable the Use Client-based Sessions option.
  - e. Click Save Changes.
3. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:
  - For `testname`, specify the value `am.AMRestAuthNSim`.
  - For `users`, specify the value `1000000` for a small cluster, `10000000` for a medium cluster, or `100000000` for a large cluster.
  - For `concurrency`, specify the value `50`.
  - For `duration`, specify the value `600`.

To run the benchmark test:

1. Install the `gatling-benchmark` Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```

2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Peak throughput	715 authNs/second	2061 authNs/second	4592 authNs/second
Peak response times	Mean: 35 ms 95th percentile: 61 ms	Mean: 29 ms 95th percentile: 48 ms	Mean: 43 ms 95th percentile: 85 ms

Metric	Small Cluster	Medium Cluster	Large Cluster
Resource requests	Memory: 20 Gi CPU: 10000 m	Memory: 8 Gi CPU: 11000 m	Memory: 4 Gi CPU: 10000 m
Resource limits	Memory: 5 Gi CPU: 3000 m	Memory: 11 Gi CPU: 12000 m	Memory: 26 Gi CPU: 16000 m
Peak pod usage	Memory: 4830 Mi CPU: 2090 m	Memory: 6327 Mi CPU: 5868 m	Memory: 19892 Mi CPU: 15072 m
Peak node usage	Memory: 14100 Mi CPU: 2290 m	Memory: 22767 Mi CPU: 10794 m	Memory: 49990 Mi CPU: 24696 m

#### 5.4.1.2. REST Login - Stateless Authentication

This benchmark test measures throughput and response times of an AM server performing stateless authentications.

Before running the benchmark test:

1. Make sure the userstore is provisioned, and the Directory Services cache is primed.  
See "To Provision the Userstores with Test Users".
2. Configure AM for stateless sessions:
  - a. Log in to the AM console as the `amadmin` user. For details, see "Accessing AM Services".
  - b. Select the top-level realm.
  - c. Select Properties.
  - d. Enable the Use Client-based Sessions option.
  - e. Click Save Changes.
3. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:
  - For `testname`, specify the value `am.AMRestAuthNSim`.
  - For `users`, specify the value `1000000` for a small cluster, `10000000` for a medium cluster, or `100000000` for a large cluster.
  - For `concurrency`, specify the value `50`.
  - For `duration`, specify the value `600`.

To run the benchmark test:

1. Install the `gatling-benchmark` Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```



2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Peak throughput	802 authNs/second	2102 authNs/second	3789 authNs/second
Peak response times	Mean: 31 ms 95th percentile: 43 ms	Mean: 25 ms 95th percentile: 53 ms	Mean: 52 ms 95th percentile: 104 ms
Resource requests	Memory: 20 Gi CPU: 10000 m	Memory: 8 Gi CPU: 11000 m	Memory: 4 Gi CPU: 10000 m
Resource limits	Memory: 5 Gi CPU: 3000 m	Memory: 11 Gi CPU: 12000 m	Memory: 26 Gi CPU: 16000 m
Peak pod usage	Memory: 4910 Mi CPU: 2100 m	Memory: 6374 Mi CPU: 6569 m	Memory: 19892 Mi CPU: 15072 m
Peak node usage	Memory: 14400 Mi CPU: 2330 m	Memory: 23206 Mi CPU: 11621 m	Memory: 49990 Mi CPU: 24696 m

### 5.4.1.3. OAuth 2.0 - Authorization Code Grant Flow - Stateful Tokens

This benchmark test measures throughput and response time of an AM server performing authentication, authorization, and session token management using stateful tokens. In this test, one transaction includes all three operations.

Before running the benchmark test:

1. Make sure the userstore is provisioned, and the Directory Services cache is primed.

See "To Provision the Userstores with Test Users".

2. Configure AM for stateful sessions:

- a. Log in to the AM console as the `amadmin` user. For details, see "Accessing AM Services".
- b. Select the top-level realm.

- c. Select Properties.
  - d. Disable the Use Client-based Sessions option.
  - e. Click Save Changes.
3. Configure AM for stateful OAuth2 tokens:
- a. Select Services > OAuth2 Provider
  - b. Disable the Use Client-based Access & Refresh Tokens option.
  - c. Click Save Changes.
4. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:
- For `testname`, specify the value `am.AMAccessTokenSim`.
  - For `users`, specify the value `1000000` for a small cluster, `10000000` for a medium cluster, or `100000000` for a large cluster.
  - For `concurrency`, specify the value `50`.
  - For `duration`, specify the value `600`.

To run the benchmark test:

1. Install the `gatling-benchmark` Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```

2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Peak throughput	774 transactions/second	2155 transactions/second	4263 transactions/second

Metric	Small Cluster	Medium Cluster	Large Cluster
Peak response times	Mean: 32 ms 95th percentile: 55 ms	Mean: 23 ms 95th percentile: 59 ms	Mean: 47 ms 95th percentile: 127 ms
Resource requests	Memory: 20 Gi CPU: 10000 m	Memory: 8 Gi CPU: 11000 m	Memory: 4 Gi CPU: 10000 m
Resource limits	Memory: 5 Gi CPU: 3000 m	Memory: 11 Gi CPU: 12000 m	Memory: 26 Gi CPU: 16000 m
Peak pod usage	Memory: 5000 Mi CPU: 2280 m	Memory: 8664 Mi CPU: 8972 m	Memory: 19892 Mi CPU: 15072 m
Peak node usage	Memory: 14200 Mi CPU: 2710 m	Memory: 21098 Mi CPU: 11072 m	Memory: 49990 Mi CPU: 24696 m

#### 5.4.1.4. OAuth 2.0 - Authorization Code Grant Flow - Stateless Tokens

This benchmark test measures throughput and response time of an AM server performing authentication, authorization, and session token management using stateless tokens. In this test, one transaction includes all three operations.

Before running the benchmark test:

1. Make sure the userstore is provisioned, and the Directory Services cache is primed.  
See "To Provision the Userstores with Test Users".
2. Configure AM for stateless sessions:
  - a. Log in to the AM console as the `amadmin` user. For details, see "Accessing AM Services".
  - b. Select the top-level realm.
  - c. Select Properties.
  - d. Enable the Use Client-based Sessions option.
  - e. Click Save Changes.
3. Configure AM for stateless OAuth2 tokens:
  - a. Select Services > OAuth2 Provider
  - b. Enable the Use Client-based Access & Refresh Tokens option.
  - c. Click Save Changes.
4. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:
  - For `testname`, specify the value `am.AMAccessTokenSim`.

- For **users**, specify the value **1000000** for a small cluster, **10000000** for a medium cluster, or **100000000** for a large cluster.
- For **concurrency**, specify the value **50**.
- For **duration**, specify the value **600**.

To run the benchmark test:

1. Install the **gatling-benchmark** Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```

2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Peak throughput	558 transactions/second	2015 transactions/second	3167 transactions/second
Peak response times	Mean: 44 ms 95th percentile: 90 ms	Mean: 26 ms 95th percentile: 74 ms	Mean: 63 ms 95th percentile: 176 ms
Resource requests	Memory: 20 Gi CPU: 10000 m	Memory: 8 Gi CPU: 11000 m	Memory: 4 Gi CPU: 10000 m
Resource limits	Memory: 5 Gi CPU: 3000 m	Memory: 11 Gi CPU: 12000 m	Memory: 26 Gi CPU: 16000 m
Peak pod usage	Memory: 5020 Mi CPU: 2360 m	Memory: 6418 Mi CPU: 8238 m	Memory: 19892 Mi CPU: 15072 m
Peak node usage	Memory: 14300 Mi CPU: 2730 m	Memory: 23225 Mi CPU: 10585 m	Memory: 49990 Mi CPU: 24696 m

## 5.4.2. Identity Management Benchmark Tests

The IDM tests examine the create, read, update, and delete (CRUD) performance through the IDM API. The Cloud Deployment Team generated CRUD load using the following simulations:

- `IDMCreateManagedUsers.scala`
- `IDMReadManagedUsers.scala`
- `IDMUpdateManagedUsers.scala`
- `IDMDeleteManagedUsers.scala`

You can specify a duration for the `IDMCreateManagedUsers.scala` simulation. The other simulations run for the required duration to complete read, update and delete the users.

### 5.4.2.1. Create

This benchmark test measures throughput and response time of an IDM server performing create operations.

Before running the benchmark test:

1. Make sure the userstore is provisioned, and the Directory Services cache is primed.  
See "To Provision the Userstores with Test Users".
2. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:
  - For `testname`, specify the value `idm.IDMCreateManagedUsers`.
  - For `users`, specify the value `1000000` for a small cluster, or `10000000` for a medium cluster.
  - For `concurrency`, specify the value `50`.
  - For `duration`, specify the value `600`.

To run the benchmark test:

1. Install the `gatling-benchmark` Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```

2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster
Number of users	1,000,000	10,000,000
Peak throughput	86 creates/second	479 creates/second
Peak response times	Mean: 173 ms 95th percentile: 266 ms	Mean: 62 ms 95th percentile: 82 ms
Resource requests	Memory: 2 Gi CPU: 1500 m	Memory: 4 Gi CPU: 2000 m
Resource limits	Memory: 3 Gi CPU: 3000 m	Memory: 6 Gi CPU: 10000 m
Peak pod usage	Memory: 4820 Mi CPU: 2210 m	Memory: 3274 Mi CPU: 8154 m
Peak node usage	Memory: 14170 Mi CPU: 2450 m	Memory: 21684 Mi CPU: 14367 m

#### 5.4.2.2. Read

This benchmark test measures throughput and response time of an IDM server performing read operations.

Before running the benchmark test

1. Make sure the userstore is provisioned, and the Directory Services cache is primed.  
See "To Provision the Userstores with Test Users".
2. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:
  - For `testname`, specify the value `idm.IDMReadManagedUsers`.
  - For `users`, specify the value `1000000` for a small cluster, or `10000000` for a medium cluster.
  - For `concurrency`, specify the value `50`.
  - For `duration`, specify the value `600`.

To run the benchmark test:

1. Install the `gatling-benchmark` Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```

2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

#### 4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster
Number of users	1,000,000	10,000,000
Peak throughput	210 reads/second	2060 reads/second
Peak response times	Mean: 141 ms 95th percentile: 223 ms	Mean: 14 ms 95th percentile: 32 ms
Resource requests	Memory: 2 Gi CPU: 1500 m	Memory: 4 Gi CPU: 2000 m
Resource limits	Memory: 3 Gi CPU: 3000 m	Memory: 6 Gi CPU: 10000 m
Peak pod usage	Memory: 4720 Mi CPU: 2010 m	Memory: 3113 Gi CPU: 6857 m
Peak node usage	Memory: 14150 Mi CPU: 2210 m	Memory: 21679 Mi CPU: 7230 m

### 5.4.2.3. Update

This benchmark test measures throughput and response time of an IDM server performing update operations.

Before running the benchmark test:

#### 1. Make sure the userstore is provisioned, and the Directory Services cache is primed.

See "To Provision the Userstores with Test Users".

#### 2. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:

- For `testname`, specify the value `idm.IDMUpdateManagedUsers`.
- For `users`, specify the value `1000000` for a small cluster, or `10000000` for a medium cluster.
- For `concurrency`, specify the value `50`.
- For `duration`, specify the value `600`.

To run the benchmark test:

1. Install the `gatling-benchmark` Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```

2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster
Number of users	1,000,000	10,000,000
Peak throughput	78 updates/second	630 updates/second
Peak response times	Mean: 384 ms 95th percentile: 615 ms	Mean: 47 ms 95th percentile: 79 ms
Resource requests	Memory: 2 Gi CPU: 1500 m	Memory: 4 Gi CPU: 2000 m
Resource limits	Memory: 3 Gi CPU: 3000 m	Memory: 6 Gi CPU: 10000 m
Peak pod usage	Memory: 2280 Mi CPU: 4920 m	Memory: 4248 Gi CPU: 9890 m
Peak node usage	Memory: 14970 Mi CPU: 2410 m	Memory: 23086 Mi CPU: 12419 m

#### 5.4.2.4. Delete

This benchmark test measures throughput and response time of an IDM server performing delete operations.

Before running the benchmark test:

1. Make sure the userstore is provisioned, and the Directory Services cache is primed.

See "To Provision the Userstores with Test Users".

2. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:

- For `testname`, specify the value `idm.IDMDeleteManagedUsers`.



- For `users`, specify the value `1000000` for a small cluster, or `10000000` for a medium cluster.
- For `concurrency`, specify the value `50`.
- For `duration`, specify the value `600`.

To run the benchmark test:

1. Install the `gatling-benchmark` Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```

2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster
Number of users	1,000,000	10,000,000
Peak throughput	102 deletes/second	853 deletes/second
Peak response times	Mean: 146 ms 95th percentile: 321 ms	Mean: 35 ms 95th percentile: 46 ms
Resource requests	Memory: 2 Gi CPU: 1500 m	Memory: 4 Gi CPU: 2000 m
Resource limits	Memory: 3 Gi CPU: 3000 m	Memory: 6 Gi CPU: 10000 m
Peak pod usage	Memory: 4820 Mi CPU: 2140 m	Memory: 4334 Gi CPU: 8824 m
Peak node usage	Memory: 14990 Mi CPU: 2360 m	Memory: 23167 Mi CPU: 11861 m

### 5.4.3. Identity Gateway Benchmark Tests

The IG tests measure throughput and response time of IG acting as a resource server and as a reverse proxy server. The Cloud Deployment Team generated load using the following simulations:

- Simulations that test IG acting as a resource server:
  - The `IGAccessTokensSim.scala` simulation, which tests a resource server that caches OAuth 2.0 tokens. With caching enabled, verification by an AM server is triggered after the first request only. Subsequent requests from the same user don't require verification.
  - The `IGAccessTokensNoCacheSim.scala` simulation, which tests a resource server with no cache. With no cache, each request must be verified by an AM server.

Both simulations require the availability of OAuth 2.0 tokens. Prior to running either simulation, you run the `IGGenerateTokensSim.scala` simulation to generate the tokens.

- The `IGReverseProxyWebSim.scala` simulation, which tests IG acting as a reverse proxy server in front of an NGINX web server. The NGINX web server returns a static web page.

### 5.4.3.1. Resource Server - Cached

This benchmark test measures throughput and response time of an IG resource server that caches tokens.

Before running the benchmark test:

1. Make sure the userstore is provisioned, and the Directory Services cache is primed.

See "To Provision the Userstores with Test Users".

2. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:

- For `testname`, specify the value `"ig.IGGenerateTokensSim ig.IGAccessTokensSim"`.
- For `oauth2_client`, specify the value `client-application`.
- For `users`, specify the value `1000000` for a small cluster, `10000000` for a medium cluster, or `100000000` for a large cluster.
- For `concurrency`, specify the value `50`.
- For `duration`, specify the value `600`.

To run the benchmark test:

1. Install the `gatling-benchmark` Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```

2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

#### 4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Peak throughput	4009 transactions/second	12666 transactions/second	5011 transactions/second
Peak response times	Mean: 6 ms 95th percentile: 18 ms	Mean: 4 ms 95th percentile: 11 ms	Mean: 10 ms 95th percentile: 16 ms
Resource requests	Memory: 2 Gi CPU: 1000 m	Memory: 2 Gi CPU: 2000 m	Memory: 2 Gi CPU: 2000 m
Resource limits	Memory: 1 Gi CPU: 1000 m	Memory: 4 Gi CPU: 4000 m	Memory: 4 Gi CPU: 4000 m
Peak pod usage	AM Memory: 1129 Mi CPU: 860 m	AM Memory: 1847 Mi CPU: 1370 m	AM Memory: 2349 Mi CPU: 3069 m

#### 5.4.3.2. Resource Server - No Cache

This benchmark test measures throughput and response time of an IG resource server with no cache. With no cache, all requests go directly to AM.

Before running the benchmark test:

##### 1. Make sure the userstore is provisioned, and the Directory Services cache is primed.

See "To Provision the Userstores with Test Users".

##### 2. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:

- For `testname`, specify the value `"ig.IGGenerateTokensSim ig.IGAccessTokensNoCacheSim"`.
- For `oauth2_client`, specify the value `client-application`.
- For `users`, specify the value `1000000` for a small cluster, `10000000` for a medium cluster, or `100000000` for a large cluster.
- For `concurrency`, specify the value `50`.
- For `duration`, specify the value `600`.

To run the benchmark test:

1. Install the `gatling-benchmark` Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```

2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Peak throughput	1719 transactions/second	4172 transactions/second	10229 transactions/second
Peak response times	Mean: 14 ms 95th percentile: 27 ms	Mean: 12 ms 95th percentile: 21 ms	Mean: 10 ms 95th percentile: 16 ms
Resource requests	Memory: 1 Gi CPU: 1000 m	Memory: 2 Gi CPU: 2000 m	Memory: 2 Gi CPU: 2000 m
Resource limits	Memory: 2 Gi CPU: 1000 m	Memory: 4 Gi CPU: 4000 m	Memory: 4 Gi CPU: 4000 m
Peak pod usage	AM Memory: 2048 Mi CPU: 870 m	AM Memory: 1707 Mi CPU: 1240 m	AM Memory: 45066 Mi CPU: 2760 m

### 5.4.3.3. Reverse Proxy

This benchmark test measures throughput and response time of an IG server as a reverse proxy in front of an nginx web server. The nginx web server returns a static web page.

Before running the benchmark test:

1. Make sure the userstore is provisioned, and the Directory Services cache is primed.

See "To Provision the Userstores with Test Users".

2. Configure the `/path/to/forgeops/helm/gatling-benchmark/values.yaml` file:

- For `testname`, specify the value `ig.IGReverseProxyWebSim`.
- For `users`, specify the value `1000000` for a small cluster, `10000000` for a medium cluster, or `100000000` for a large cluster.
- For `concurrency`, specify the value `50`.
- For `duration`, specify the value `600`.

To run the benchmark test:

1. Install the `gatling-benchmark` Helm chart:

```
$ helm install --name benchmark helm/gatling-benchmark
```

2. Send the logs to the terminal window:

```
$ ./get-logs.sh
```

3. (Optional) View the results in the Gatling UI.

See "To View Test Results in a Browser".

4. (Optional) Stop the benchmark test:

```
$ helm del --purge benchmark
```

See also "Running AM, IDM, or IG Benchmark Tests".

You can compare your results with the Cloud Deployment Team results here:

Metric	Small Cluster	Medium Cluster	Large Cluster
Number of users	1,000,000	10,000,000	100,000,000
Peak throughput	4055 transactions/second	12942 transactions/second	12489 transactions/second
Peak response times	Mean: 6 ms 95th percentile: 18 ms	Mean: 4 ms 95th percentile: 10 ms	Mean: 8 ms 95th percentile: 22 ms
Resource requests	Memory: 2 Gi CPU: 1000 m	Memory: 2 Gi CPU: 2000 m	Memory: 2 Gi CPU: 2000 m
Resource limits	Memory: 2 Gi CPU: 1000 m	Memory: 4 Gi CPU: 4000 m	Memory: 4 Gi CPU: 4000 m
Peak pod usage	AM Memory: 1019 Mi CPU: 1003 m	AM Memory: 1702 Mi CPU: 2900 m	AM Memory: 1744 Mi CPU: 2431 m

## Chapter 6

# Taking the Next Steps

If you've followed the instructions in this cookbook *without modifying configurations*, then the following indicate that you've successfully deployed the CDM:

- The Kubernetes cluster and pods are up and running.
- DS, AM, IDM, and IG are installed and running. You can access each ForgeRock component.
- DS is provisioned with users. Replication and failover work as expected.
- Monitoring tools are installed and running. You can access a monitoring console for DS, AM, IDM, and IG.
- You can run the benchmarking tests in this cookbook without errors.
- Your benchmarking test results are similar to the benchmarks reported in this cookbook.

When you're satisfied that all of these conditions are met, then you've completed your initial deployment phase. Congratulations!

Now you're ready to engineer your ForgeRock Identity Platform site reliability. See the [ForgeRock Site Reliability Guides](#) for information about:

- Modifying the CDM configuration to suit your business needs.
- Monitoring site health and performance.
- Backing up configuration and user data for disaster preparedness.
- Securing your site.

The CDM and its documentation continue to evolve. As they evolve, the *ForgeRock Site Reliability Guides* will become your resources for information about CDM continuous delivery and continuous integration.

# Appendix A. Getting Support

This appendix contains information about support options for the ForgeRock DevOps Examples and the ForgeRock Identity Platform.

## A.1. ForgeRock DevOps Support

ForgeRock has developed artifacts in the `forgeops` and `forgeops-init` Git repositories for the purpose of deploying the ForgeRock Identity Platform in the cloud. The companion ForgeRock DevOps documentation provides examples, including the ForgeRock Cloud Deployment Model (CDM), to help you get started.

These artifacts and documentation are provided on an "as is" basis. ForgeRock does not guarantee the individual success developers may have in implementing the code on their development platforms or in production configurations.

### A.1.1. Commercial Support

ForgeRock provides commercial support for the following DevOps resources:

- Dockerfiles and Helm charts in the `forgeops` Git repository
- ForgeRock [DevOps guides](#).

ForgeRock provides commercial support for the ForgeRock Identity Platform. For supported components, containers, and Java versions, see the following:

- *ForgeRock Access Management Release Notes*
- *ForgeRock Identity Management Release Notes*

- [ForgeRock Directory Services Release Notes](#)
- [ForgeRock Identity Message Broker Release Notes](#)
- [ForgeRock Identity Gateway Release Notes](#)

### A.1.2. Support Limitations

ForgeRock provides no commercial support for the following:

- Artifacts other than Dockerfiles or Helm charts in the `forgeops` and `forgeops-init` repositories. Examples include scripts, example configurations, and so forth.
- Non-ForgeRock infrastructure. Examples include Docker, Kubernetes, Google Cloud Platform, Amazon Web Services, and so forth.
- Non-ForgeRock software. Examples include Java, Apache Tomcat, NGINX, Apache HTTP Server, and so forth.
- Production deployments that use the DevOps evaluation-only Docker images. When deploying the ForgeRock Identity Platform using Docker images, you must build and use your own images for production deployments. For information about how to build Docker images for the ForgeRock Identity Platform, see *"Building and Pushing Docker Images"* in the *DevOps Developer's Guide*.

### A.1.3. Third-Party Kubernetes Services

ForgeRock supports deployments on Google Kubernetes Engine (GKE), Amazon Elastic Kubernetes Service (Amazon EKS), Microsoft Azure Kubernetes Service (AKS), and Red Hat OpenShift.

Red Hat OpenShift is a tested and supported platform using Kubernetes for deployment. ForgeRock uses OpenShift tools such as Minishift, as well as other representative environments such as Amazon AWS for the testing. We do not test using bare metal due to the many customer permutations of deployment and configuration that may exist, and therefore cannot guarantee that we have tested in the same way a customer chooses to deploy. We will make commercially reasonable efforts to provide first-line support for any reported issue. In the case we are unable to reproduce a reported issue internally, we will request the customer engage OpenShift support to collaborate on problem identification and remediation. Customers deploying on OpenShift are expected to have a support contract in place with IBM/Red Hat that ensures support resources can be engaged if this situation may occur.

## A.2. Accessing Documentation Online

ForgeRock publishes comprehensive documentation online:

- The ForgeRock [Knowledge Base](#) offers a large and increasing number of up-to-date, practical articles that help you deploy and manage ForgeRock software.



While many articles are visible to community members, ForgeRock customers have access to much more, including advanced information for customers using ForgeRock software in a mission-critical capacity.

- ForgeRock product documentation, such as this document, aims to be technically accurate and complete with respect to the software documented. It is visible to everyone and covers all product features and examples of how to use them.

## A.3. How to Report Problems or Provide Feedback

If you are a named customer Support Contact, contact ForgeRock using the [Customer Support Portal](#) to request information or report a problem with Dockerfiles or Helm charts in the DevOps Examples or the CDM.

If you have questions regarding the DevOps Examples or the CDM that are not answered in the documentation, file an issue at <https://github.com/ForgeRock/forgeops/issues>.

When requesting help with a problem, include the following information:

- Description of the problem, including when the problem occurs and its impact on your operation.
- Steps to reproduce the problem.

If the problem occurs on a Kubernetes system other than Minikube, GKE, EKS, OpenShift, or AKS, we might ask you to reproduce the problem on one of those.

- HTML output from the **debug-logs.sh** script. For more information, see "Running the debug-logs.sh Script" in the *DevOps Developer's Guide*.
- Description of the environment, including the following information:
  - Environment type: Minikube, GKE, EKS, AKS, or OpenShift.
  - Software versions of supporting components:
    - Oracle VirtualBox (Minikube environments only).
    - Docker client (all environments).
    - Minikube (all environments).
    - **kubectrl** command (all environments).
    - Kubernetes Helm (all environments).
    - Google Cloud SDK (GKE environments only).
    - Amazon AWS Command Line Interface (EKS environments only).

- Azure Command Line Interface (AKS environments only).
- **forgeops** repository branch.
- Any patches or other software that might be affecting the problem.

## A.4. Getting Support and Contacting ForgeRock

ForgeRock provides support services, professional services, training through ForgeRock University, and partner services to assist you in setting up and maintaining your deployments. For a general overview of these services, see <https://www.forgerock.com>.

ForgeRock has staff members around the globe who support our international customers and partners. For details on ForgeRock's support offering, including support plans and service level agreements (SLAs), visit <https://www.forgerock.com/support>.

# Glossary

affinity (AM)	<p>AM affinity based load balancing ensures that the CTS token creation load is spread over multiple server instances (the token origin servers). Once a CTS token is created and assigned to a session, all subsequent token operations are sent to the same token origin server from any AM node. This ensures that the load of CTS token management is spread across directory servers.</p> <p>Source: <i>Best practices for using Core Token Service (CTS) Affinity based load balancing in AM</i></p>
Amazon EKS	<p>Amazon Elastic Container Service for Kubernetes (Amazon EKS) is a managed service that makes it easy for you to run Kubernetes on Amazon Web Services without needing to set up or maintain your own Kubernetes control plane.</p> <p>Source: <i>What is Amazon EKS</i> in the Amazon EKS documentation.</p>
ARN (AWS)	<p>An Amazon Resource Name (ARN) uniquely identifies an Amazon Web Service (AWS) resource. AWS requires an ARN when you need to specify a resource unambiguously across all of AWS, such as in IAM policies and API calls.</p> <p>Source: <i>Amazon Resource Names (ARNs) and AWS Service Namespaces</i> in the AWS documentation.</p>
AWS IAM Authenticator for Kubernetes	<p>The AWS IAM Authenticator for Kubernetes is an authentication tool that enables you to use <i>Amazon Web Services (AWS)</i> credentials for authenticating to a Kubernetes cluster.</p> <p>Source: <i>AWS IAM Authenticator for Kubernetes</i> <a href="#">README</a> file on <a href="#">GitHub</a>.</p>

cloud-controller-manager	<p>The <code>cloud-controller-manager</code> daemon runs controllers that interact with the underlying cloud providers. <code>cloud-controller-manager</code> is an alpha feature introduced in Kubernetes release 1.6. The <code>cloud-controller-manager</code> daemon runs cloud-provider-specific controller loops only.</p> <p>Source: <i>cloud-controller-manager</i> section in the Kubernetes Concepts documentation.</p>
Cloud Developer's Kit (CDK)	<p>The developer artifacts in the <code>forgeops</code> Git repository, together with the ForgeRock Identity Platform documentation form the Cloud Developer's Kit (CDK). Use the CDK to stand up the platform in your developer environment.</p>
Cloud Deployment Model (CDM)	<p>The Cloud Deployment Model (CDM) is a common use ForgeRock Identity Platform architecture, designed to be easy to deploy and easy to replicate. The ForgeRock Cloud Deployment Team has developed Helm charts, Docker images, and other artifacts expressly to build the CDM.</p>
CloudFormation (AWS)	<p>CloudFormation is a service that helps you model and set up your Amazon Web Services (AWS) resources. You create a template that describes all the AWS resources that you want. AWS CloudFormation takes care of provisioning and configuring those resources for you.</p> <p>Source: <i>What is AWS CloudFormation?</i> in the AWS documentation.</p>
CloudFormation template (AWS)	<p>An AWS CloudFormation template describes the resources that you want to provision in your AWS stack. AWS CloudFormation templates are text files formatted in JSON or YAML.</p> <p>Source: <i>Working with AWS CloudFormation Templates</i> in the AWS documentation.</p>
cluster	<p>A container cluster is the foundation of Kubernetes Engine. A cluster consists of at least one <code>cluster master</code> and multiple worker machines called nodes. The Kubernetes objects that represent your containerized applications all run on top of a cluster.</p> <p>Source: <i>Container Cluster Architecture</i> in the Kubernetes Concepts documentation.</p>
cluster master	<p>A cluster master schedules, runs, scales and upgrades the workloads on all nodes of the cluster. The cluster master also manages network and storage resources for workloads.</p> <p>Source: <i>Container Cluster Architecture</i> in the Kubernetes Concepts documentation.</p>
ConfigMap	<p>A configuration map, called <code>ConfigMap</code> in Kubernetes manifests, binds the configuration files, command-line arguments, environment</p>

variables, port numbers, and other configuration artifacts to the assigned containers and system components at runtime. The configuration maps are useful for storing and sharing non-sensitive, unencrypted configuration information.

Source: *ConfigMap* in the [Kubernetes Concepts](#) documentation.

#### container

A container is an allocation of resources such as CPU, network I/O, bandwidth, block I/O, and memory that can be “contained” together and made available to specific processes without interference from the rest of the system.

Source *Container Cluster Architecture* in the [Google Cloud Platform](#) documentation

#### DaemonSet

A set of daemons, called **DaemonSet** in Kubernetes manifests, manages a group of replicated pods. Usually, the daemon set follows an one-pod-per-node model. As you add nodes to a node pool, the daemon set automatically distributes the pod workload to the new nodes as needed.

Source *DaemonSet* in the [Google Cloud Platform](#) documentation.

#### Deployment

A Kubernetes deployment represents a set of multiple, identical pods. A Kubernetes deployment runs multiple replicas of your application and automatically replaces any instances that fail or become unresponsive.

Source: *Deployment* in the [Google Cloud Platform](#) documentation.

#### deployment controller

A deployment controller provides declarative updates for pods and replica sets. You describe a desired state in a deployment object, and the deployment controller changes the actual state to the desired state at a controlled rate. You can define deployments to create new replica sets, or to remove existing deployments and adopt all their resources with new deployments.

Source: *Deployments* in the [Google Cloud Platform](#) documentation.

#### Docker Cloud

Docker Cloud provides a hosted registry service with build and testing facilities for Dockerized application images; tools to help you set up and manage host infrastructure; and application lifecycle features to automate deploying (and redeploying) services created from images.

Source: *About Docker Cloud* in the [Docker Cloud](#) documentation.

#### Docker container

A Docker container is a runtime instance of a [Docker image](#). A Docker container is isolated from other containers and its host machine. You can control how isolated your container’s network,

storage, or other underlying subsystems are from other containers or from the host machine.

Source: Containers section in the Docker architecture documentation.

#### Docker daemon

The Docker daemon (`dockerd`) listens for Docker API requests and manages Docker objects such as images, containers, networks, and volumes. A Docker daemon can also communicate with other Docker daemons to manage Docker services.

Source: *Docker daemon* section in the Docker Overview documentation.

#### Docker Engine

The Docker Engine is a client-server application with these components:

- A server, which is a type of long-running program called a daemon process (the `dockerd` command)
- A REST API, which specifies interfaces that programs can use to talk to the daemon and tell it what to do
- A command-line interface (CLI) client (the `docker` command)

Source: Docker Engine section in the Docker Overview documentation.

#### Dockerfile

A Dockerfile is a text file that contains the instructions for building a Docker image. Docker uses the Dockerfile to automate the process of building a Docker image.

Source: *Dockerfile* section in the Docker Overview documentation.

#### Docker Hub

Docker Hub provides a place for you and your team to build and ship Docker images. You can create public repositories that can be accessed by any other Docker Hub user, or you can create private repositories you can control access to.

An image is an application you would like to run. A container is a running instance of an image.

Source: *Overview of Docker Hub* section in the Docker Overview documentation.

#### Docker image

A Docker image is a read-only template with instructions for creating a Docker container. Often, an image is based on another image, with some additional customization.

A Docker image includes the application code, a runtime engine, libraries, environment variables, and configuration files that are required to run the application.

An image is an application you would like to run. A container is a running instance of an image.

Source: *Docker objects* section in the Docker Overview documentation. [Hello Whales: Images vs. Containers in Dockers](#).

#### Docker namespace

Docker namespaces provide a layer of isolation. When you run a container, Docker creates a set of namespaces for that container. Each aspect of a container runs in a separate namespace and its access is limited to that namespace.

The **PID** namespace is the mechanism for remapping process IDs inside the container. Other namespaces such as net, mnt, ipc, and uts provide the isolated environments we know as containers. The user namespace is the mechanism for remapping user IDs inside a container.

Source: *Namespaces* section in the Docker Overview documentation.

#### Docker registry

A Docker registry stores [Docker images](#). Docker Hub and Docker Cloud are public registries that anyone can use, and Docker is configured to look for images on [Docker Hub](#) by default. You can also run your own private registry.

Source: *Docker registries* section in the Docker Overview documentation.

#### Docker repository

A Docker repository is a public, certified repository from vendors and contributors to Docker. It contains [Docker images](#) that you can use as the foundation to build your applications and services.

Source: *Repositories on Docker Hub* section in the Docker Overview documentation.

#### Docker service

In a distributed application, different pieces of the application are called “services.” Docker services are really just “containers in production.” A Docker service runs only one image, but it codifies the way that image runs including which ports to use, the number replicas the container should run, and so on. By default, the services are load-balanced across all worker nodes.

Source: *About services* in the Docker Get Started documentation.

#### dynamic volume provisioning

The process of creating storage volumes on demand is called dynamic volume provisioning. Dynamic volume provisioning allows storage

volumes to be created on-demand. It automatically provisions storage when it is requested by users.

Source: *Dynamic Volume Provisioning* in the Kubernetes Concepts documentation.

egress

An egress controls access to destinations outside the network from within a Kubernetes network. For an external destination to be accessed from a Kubernetes environment, the destination should be listed as an allowed destination in the whitelist configuration.

Source: *Network Policies* in the Kubernetes Concepts documentation.

firewall rule

A firewall rule lets you allow or deny traffic to and from your virtual machine instances based on a configuration you specify. Each Kubernetes network has a set of firewall rules controlling access to and from instances in its subnets. Each firewall rule is defined to apply to either incoming [glossary-ingress](#)(ingress) or outgoing (egress) traffic, not both.

Source: *Firewall Rules Overview* in the Google Cloud Platform documentation.

garbage collection

Garbage collection is the process of deleting unused objects. [Kubelets](#) perform garbage collection for containers every minute and garbage collection for images every five minutes. You can adjust the high and low threshold flags and garbage collection policy to tune image garbage collection.

Source: *Garbage Collection* in the Kubernetes Concepts documentation.

Google Kubernetes Engine (GKE)

The Google Kubernetes Engine (GKE) is an environment for deploying, managing, and scaling your containerized applications using Google infrastructure. The GKE environment consists of multiple machine instances grouped together to form a [container cluster](#).

Source: *Kubernetes Engine Overview* in the Google Cloud Platform documentation.

ingress

An ingress is a collection of rules that allow inbound connections to reach the cluster services.

Source: *Ingress* in the Kubernetes Concepts documentation.

instance group

An instance group is a collection of instances of virtual machines. The instance groups enable you to easily monitor and control the group of virtual machines together.



	<p>Source: <i>Instance Groups</i> in the Google Cloud Platform documentation.</p>
instance template	<p>An instance template is a global API resource that you can use to create VM instances and managed instance groups. Instance templates define the machine type, image, zone, labels, and other instance properties. They are very helpful in replicating the environments.</p> <p>Source: <i>Instance Templates</i> in the Google Cloud Platform documentation.</p>
kubectrl	<p>The <b>kubectrl</b> command-line tool supports several different ways to create and manage Kubernetes objects.</p> <p>Source: <i>Kubernetes Object Management</i> in the Kubernetes Concepts documentation.</p>
kube-controller-manager	<p>The Kubernetes controller manager is a process that embeds core controllers that are shipped with Kubernetes. Logically each controller is a separate process, but to reduce complexity, they are all compiled into a single binary and run in a single process.</p> <p>Source: <i>kube-controller-manager</i> in the Kubernetes Reference documentation.</p>
kubelet	<p>A kubelet is an agent that runs on each node in the cluster. It ensures that containers are running in a pod.</p> <p>Source: <i>kubelets</i> in the Kubernetes Concepts documentation.</p>
kube-scheduler	<p>The <b>kube-scheduler</b> component is on the master node and watches for newly created pods that do not have a node assigned to them, and selects a node for them to run on.</p> <p>Source: <i>Kubernetes components</i> in the Kubernetes Concepts documentation.</p>
Kubernetes	<p>Kubernetes is an open source platform designed to automate deploying, scaling, and operating application containers.</p> <p>Source: <i>Kubernetes Concepts</i></p>
Kubernetes DNS	<p>A Kubernetes DNS pod is a pod used by the kubelets and the individual containers to resolve DNS names in the cluster.</p> <p>Source: <i>DNS for services and pods</i> in the Kubernetes Concepts documentation.</p>

Kubernetes namespace	<p>A Kubernetes namespace is a virtual cluster that provides a way to divide cluster resources between multiple users. Kubernetes starts with three initial namespaces:</p> <ul style="list-style-type: none"><li>• <b>default</b>: The default namespace for user created objects which don't have a namespace</li><li>• <b>kube-system</b>: The namespace for objects created by the Kubernetes system</li><li>• <b>kube-public</b>: The automatically created namespace that is readable by all users</li></ul> <p>Kubernetes supports multiple virtual clusters backed by the same physical cluster.</p> <p>Source: <i>Namespaces</i> in the Kubernetes Concepts documentation.</p>
Let's Encrypt	<p>Let's Encrypt is a free, automated, and open certificate authority.</p> <p>Source: Let's Encrypt web site.</p>
network policy	<p>A Kubernetes network policy specifies how groups of pods are allowed to communicate with each other and with other network endpoints.</p> <p>Source: <i>Network policies</i> in the Kubernetes Concepts documentation.</p>
node (Kubernetes)	<p>A Kubernetes node is a virtual or physical machine in the cluster. Each node is managed by the master components and includes the services needed to run the pods.</p> <p>Source: <i>Nodes</i> in the Kubernetes Concepts documentation.</p>
node controller (Kubernetes)	<p>A Kubernetes node controller is a Kubernetes master component that manages various aspects of the nodes such as: lifecycle operations on the nodes, operational status of the nodes, and maintaining an internal list of nodes.</p> <p>Source: <i>Node Controller</i> in the Kubernetes Concepts documentation.</p>
persistent volume	<p>A persistent volume (PV) is a piece of storage in the cluster that has been provisioned by an administrator. It is a resource in the cluster just like a node is a cluster resource. PVs are volume plugins that have a lifecycle independent of any individual pod that uses the PV.</p> <p>Source: <i>Persistent Volumes</i> in the Kubernetes Concepts documentation.</p>
persistent volume claim	<p>A persistent volume claim (PVC) is a request for storage by a user. A PVC specifies size, and access modes such as:</p>

- Mounted once for read and write access
- Mounted many times for read-only access

Source: *Persistent Volumes* in the Kubernetes Concepts documentation.

pod anti-affinity  
(Kubernetes)

Kubernetes pod anti-affinity allows you to constrain which nodes can run your pod, based on labels on the **pods** that are already running on the node rather than based on labels on nodes. Pod anti-affinity enables you to control the spread of workload across nodes and also isolate failures to nodes.

Source: *Inter-pod affinity and anti-affinity*

pod (Kubernetes)

A Kubernetes pod is the smallest, most basic deployable object in Kubernetes. A pod represents a single instance of a running process in a cluster. Containers within a pod share an IP address and port space.

Source: *Understanding Pods* in the Kubernetes Concepts documentation.

replication controller

A replication controller ensures that a specified number of Kubernetes pod replicas are running at any one time. The **replication controller** ensures that a pod or a homogeneous set of pods is always up and available.

Source: *ReplicationController* in the Kubernetes Concepts documentation.

secret (Kubernetes)

A Kubernetes secret is a secure object that stores sensitive data, such as passwords, OAuth 2.0 tokens, and SSH keys in your clusters.

Source: *Secrets* in the Kubernetes Concepts documentation.

security group (AWS)

A security group acts as a virtual firewall that controls the traffic for one or more compute instances.

Source: *Amazon EC2 Security Groups* in the AWS documentation.

service (Kubernetes)

A Kubernetes service is an abstraction which defines a logical set of pods and a policy by which to access them. This is sometimes called a microservice.

Source: *Services* in the Kubernetes Concepts documentation.

shard

Sharding is a way of partitioning directory data so that the load can be shared by multiple directory servers. Each data partition, also

known as a *shard*, exposes the same set of naming contexts, but only a subset of the data. For example, a distribution might have two shards. The first shard contains all users whose name begins with A-M, and the second contains all users whose name begins with N-Z. Both have the same naming context.

Source: *Class Partition* in the *OpenDJ Javadoc*.

stack (AWS)

A stack is a collection of AWS resources that you can manage as a single unit. You can create, update, or delete a collection of resources by using stacks. All the resources in a stack are defined by the [template](#).

Source: *Working with Stacks* in the AWS documentation.

stack set (AWS)

A stack set is a container for stacks. You can provision stacks across AWS accounts and regions by using a single AWS [template](#). All the resources included in each stack of a stack set are defined by the same template.

Source: *StackSets Concepts* in the AWS documentation.

volume (Kubernetes)

A Kubernetes volume is a storage volume that has the same lifetime as the pod that encloses it. Consequently, a volume outlives any containers that run within the pod, and data is preserved across container restarts. When a pod ceases to exist, the Kubernetes volume also ceases to exist.

Source: *Volumes* in the Kubernetes Concepts documentation.

VPC (AWS)

A virtual private cloud (VPC) is a virtual network dedicated to your AWS account. It is logically isolated from other virtual networks in the AWS Cloud.

Source: *What Is Amazon VPC?* in the AWS documentation.

worker node (AWS)

An Amazon Elastic Container Service for Kubernetes (Amazon EKS) worker node is a standard compute instance provisioned in Amazon EKS.

Source: *Worker Nodes* in the AWS documentation.

workload (Kubernetes)

A Kubernetes workload is the collection of applications and batch jobs packaged into a [container](#). Before you deploy a workload on a cluster, you must first package the workload into a container.

Source: *Understanding Pods* in the Kubernetes Concepts documentation.