



Getting Started on Minishift

/ ForgeRock Identity Platform 6.5

Latest update: 6.5.2

Gina Cariaga
David Goldsmith
Shankar Raman

ForgeRock AS.
201 Mission St., Suite 2900
San Francisco, CA 94105, USA
+1 415-599-1100 (US)
www.forgerock.com

Copyright © 2019 ForgeRock AS.

Abstract

Quick introduction to ForgeRock Identity Platform™ deployment on Minishift for new users and readers evaluating the software.



This work is licensed under the Creative Commons Attribution-NonCommercial-NoDerivs 3.0 Unported License.

To view a copy of this license, visit <https://creativecommons.org/licenses/by-nc-nd/3.0/> or send a letter to Creative Commons, 444 Castro Street, Suite 900, Mountain View, California, 94041, USA.

ForgeRock® and ForgeRock Identity Platform™ are trademarks of ForgeRock Inc. or its subsidiaries in the U.S. and in other countries. Trademarks are the property of their respective owners.

UNLESS OTHERWISE MUTUALLY AGREED BY THE PARTIES IN WRITING, LICENSOR OFFERS THE WORK AS-IS AND MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND CONCERNING THE WORK, EXPRESS, IMPLIED, STATUTORY OR OTHERWISE, INCLUDING, WITHOUT LIMITATION, WARRANTIES OF TITLE, MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NONINFRINGEMENT, OR THE ABSENCE OF LATENT OR OTHER DEFECTS, ACCURACY, OR THE PRESENCE OF ABSENCE OF ERRORS, WHETHER OR NOT DISCOVERABLE. SOME JURISDICTIONS DO NOT ALLOW THE EXCLUSION OF IMPLIED WARRANTIES, SO SUCH EXCLUSION MAY NOT APPLY TO YOU.

EXCEPT TO THE EXTENT REQUIRED BY APPLICABLE LAW, IN NO EVENT WILL LICENSOR BE LIABLE TO YOU ON ANY LEGAL THEORY FOR ANY SPECIAL, INCIDENTAL, CONSEQUENTIAL, PUNITIVE OR EXEMPLARY DAMAGES ARISING OUT OF THIS LICENSE OR THE USE OF THE WORK, EVEN IF LICENSOR HAS BEEN ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.

DejaVu Fonts

Bitstream Vera Fonts Copyright

Copyright (c) 2003 by Bitstream, Inc. All Rights Reserved. Bitstream Vera is a trademark of Bitstream, Inc.

Permission is hereby granted, free of charge, to any person obtaining a copy of the fonts accompanying this license ("Fonts") and associated documentation files (the "Font Software"), to reproduce and distribute the Font Software, including without limitation the rights to use, copy, merge, publish, distribute, and/or sell copies of the Font Software, and to permit persons to whom the Font Software is furnished to do so, subject to the following conditions:

The above copyright and trademark notices and this permission notice shall be included in all copies of one or more of the Font Software typefaces.

The Font Software may be modified, altered, or added to, and in particular the designs of glyphs or characters in the Fonts may be modified and additional glyphs or characters may be added to the Fonts, only if the fonts are renamed to names not containing either the words "Bitstream" or the word "Vera".

This License becomes null and void to the extent applicable to Fonts or Font Software that has been modified and is distributed under the "Bitstream Vera" names.

The Font Software may be sold as part of a larger software package but no copy of one or more of the Font Software typefaces may be sold by itself.

THE FONT SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT OF COPYRIGHT, PATENT, TRADEMARK, OR OTHER RIGHT. IN NO EVENT SHALL BITSTREAM OR THE GNOME FOUNDATION BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, INCLUDING ANY GENERAL, SPECIAL, INDIRECT, INCIDENTAL, OR CONSEQUENTIAL DAMAGES, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF THE USE OR INABILITY TO USE THE FONT SOFTWARE OR FROM OTHER DEALINGS IN THE FONT SOFTWARE.

Except as contained in this notice, the names of Gnome, the Gnome Foundation, and Bitstream Inc., shall not be used in advertising or otherwise to promote the sale, use or other dealings in this Font Software without prior written authorization from the Gnome Foundation or Bitstream Inc., respectively. For further information, contact: fonts at gnome dot org.

Arev Fonts Copyright

Copyright (c) 2006 by Tavmjong Bah. All Rights Reserved.

Permission is hereby granted, free of charge, to any person obtaining a copy of the fonts accompanying this license ("Fonts") and associated documentation files (the "Font Software"), to reproduce and distribute the modifications to the Bitstream Vera Font Software, including without limitation the rights to use, copy, merge, publish, distribute, and/or sell copies of the Font Software, and to permit persons to whom the Font Software is furnished to do so, subject to the following conditions:

The above copyright and trademark notices and this permission notice shall be included in all copies of one or more of the Font Software typefaces.

The Font Software may be modified, altered, or added to, and in particular the designs of glyphs or characters in the Fonts may be modified and additional glyphs or characters may be added to the Fonts, only if the fonts are renamed to names not containing either the words "Tavmjong Bah" or the word "Arev".

This License becomes null and void to the extent applicable to Fonts or Font Software that has been modified and is distributed under the "Tavmjong Bah Arev" names.

The Font Software may be sold as part of a larger software package but no copy of one or more of the Font Software typefaces may be sold by itself.

THE FONT SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT OF COPYRIGHT, PATENT, TRADEMARK, OR OTHER RIGHT. IN NO EVENT SHALL TAVMJONG BAH BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, INCLUDING ANY GENERAL, SPECIAL, INDIRECT, INCIDENTAL, OR CONSEQUENTIAL DAMAGES, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF THE USE OR INABILITY TO USE THE FONT SOFTWARE OR FROM OTHER DEALINGS IN THE FONT SOFTWARE.

Except as contained in this notice, the name of Tavmjong Bah shall not be used in advertising or otherwise to promote the sale, use or other dealings in this Font Software without prior written authorization from Tavmjong Bah. For further information, contact: tavmjong @ free . fr.

FontAwesome Copyright

Copyright (c) 2017 by Dave Gandy, <http://fontawesome.io>.

This Font Software is licensed under the SIL Open Font License, Version 1.1. See <https://opensource.org/licenses/OFL-1.1>.

Table of Contents

Preface	iv
1. Introducing ForgeRock Identity Platform on Minishift	1
1.1. About the Example Deployment	1
2. Setting Up the Deployment Environment	3
2.1. Installing Required Third-Party Software	3
2.2. Cloning the forgeops Repository	3
2.3. Initiating Minishift	4
3. Deploying AM and IDM	5
3.1. Deploying Access Management	5
3.2. Creating Routes	6
3.3. Verifying AM Access	7
3.4. Deploying Identity Management	7
3.5. Verifying IDM Access	8
A. Getting Support	9
A.1. ForgeRock DevOps Support	9
A.2. Accessing Documentation Online	10
A.3. How to Report Problems or Provide Feedback	11
A.4. Getting Support and Contacting ForgeRock	12
Glossary	13

Preface

The *Getting Started on Minishift* guide provides instructions for quickly installing AM and IDM on a Minishift environment.

This guide covers the tasks you need to quickly get the AM and IDM on a Minishift environment.

Before You Begin

Before deploying the ForgeRock Identity Platform in a DevOps environment, read the important information in [Start Here](#).

About ForgeRock Identity Platform Software

ForgeRock Identity Platform™ serves as the basis for our simple and comprehensive Identity and Access Management solution. We help our customers deepen their relationships with their customers, and improve the productivity and connectivity of their employees and partners. For more information about ForgeRock and about the platform, see <https://www.forgerock.com>.

The platform includes the following components:

- ForgeRock® Access Management (AM)
- ForgeRock® Identity Management (IDM)
- ForgeRock® Directory Services (DS)
- ForgeRock® Identity Gateway (IG)

Chapter 1

Introducing ForgeRock Identity Platform on Minishift

Minishift is a single-node OpenShift cluster environment running in a local virtual machine.

This *Getting Started on Minishift* guide provides instructions for quickly deploying and running the ForgeRock Identity Platform on a Minishift environment.

1.1. About the Example Deployment

The example deployment presented in this guide lets you get a simple ForgeRock Identity Platform deployment up and running in Minishift as quickly as possible. The deployment uses a minimally viable configuration for AM and IDM. This minimal configuration is suitable for evaluation and demonstration purposes only.

This section describes several characteristics of the example deployment, and provides resources you can use for more complex deployments.

1.1.1. ForgeRock Identity Platform Configuration

The example deployment configures ForgeRock Identity Platform components as simply as possible:

- AM's configuration is empty: no realms, service configurations, or policies are configured in addition to the default configuration.
- IDM's configuration implements bidirectional data synchronization between IDM and LDAP described in [Synchronizing Data Between LDAP and IDM](#) in the *Samples Guide*.

1.1.2. Secure Communication With ForgeRock Identity Platform Services

The example deployment provides secure access over HTTPS to ForgeRock Identity Platform server web UIs and REST APIs.

See "Configuring and Installing the frconfig Helm Chart" in the *DevOps Developer's Guide* for more information about securing communication to ForgeRock Identity Platform servers.

1.1.3. Runtime Changes to the AM Web Application

The example deployment installs the default AM `.war` file. You can customize this `.war` file to provide enhancements such as custom authentication modules, cross-origin resource sharing (CORS) support, or a custom look and feel for web UIs.

See "Customizing the AM Web Application" in the *DevOps Developer's Guide* for details about customizing the AM `.war` file when running in a DevOps environment.

Chapter 2

Setting Up the Deployment Environment

This chapter describes how to set up your Minishift environment to deploy AM and IDM.

The chapter covers the following topics:

- "Installing Required Third-Party Software"
- "Cloning the forgeops Repository"
- "Initiating Minishift"

2.1. Installing Required Third-Party Software

Before installing the AM and IDM, review "Installing Required Third-Party Software" in the *DevOps Release Notes* to determine which software you need. Then install the required software on your local computer.

2.2. Cloning the forgeops Repository

Before you can deploy the AM environment, you must clone the `forgeops` repository. Follow the instructions in "forgeops Repository" in the *DevOps Release Notes*.

To Obtain the forgeops Repository

The `forgeops` repository is a public Git repository. You do not need credentials to clone it:

1. Clone the `forgeops` repository:

```
$ git clone https://github.com/ForgeRock/forgeops.git
```

2. Check out the `release/6.5.2` branch:

```
$ cd forgeops  
$ git checkout release/6.5.2
```

2.3. Initiating Minishift

This section describes how to start up Minishift and enable appropriate security context constraints (SCCs). SCCs allow administrators to control permissions for pods.

To Set Up Minishift

1. Start your Minishift instance:

```
$ minishift start
```

2. Enable the **anyuid** policy. This policy allows pods to run as the forgerock user, instead of a random user id generated by Minishift:

```
$ minishift addons enable anyuid
```

3. Log in to Minishift as the system user and get the available SCCs:

```
$ oc login -u system:admin  
$ oc get scc
```

4. Log back in to Minishift as the normal user:

```
$ oc login -u developer -p password
```


Chapter 3

Deploying AM and IDM

This chapter describes how to deploy AM and IDM in a Minishift environment.

This chapter covers the following topics:

- "Deploying Access Management"
- "Creating Routes"
- "Verifying AM Access"
- "Deploying Identity Management"
- "Verifying IDM Access "

3.1. Deploying Access Management

Perform the following procedure to install, configure, and start AM:

To Deploy Access Management

1. Change to the Helm charts directory in your `forgeops` repository clone:

```
$ cd /path/to/forgeops/helm
```

2. Deploy the `frconfig` Helm chart:

```
$ helm template frconfig | oc apply -f -
```

3. Deploy the configuration store:

```
$ helm template --set instance=configstore ds | oc apply -f -
```

4. Deploy the user store:

```
$ helm template --set instance=userstore ds | oc apply -f -
```

5. Deploy Amster:

```
$ helm template amster | oc apply -f -
```

6. Deploy AM:

```
$ helm template openam | oc apply -f -
```

After completing the "To Deploy Access Management" procedure, AM is deployed in the `myproject` namespace of Minishift.

3.2. Creating Routes

A route is the Minishift equivalent of a Kubernetes ingress. Create routes that let users access the AM console, the IDM user interface, and the IDM admin console.

To Create Routes

1. Access the Minishift console:

```
$ minishift console
```

2. Note the Minishift IP address, for example: 192.168.99.102.

3. Add host entries to your `/etc/hosts` file.

For example:

```
192.168.99.102 openam login.myproject.iam.example.com myproject.iam.example.com
```

4. In the Minishift console, create three routes to provide access to AM and IDM UIs:

a. Create a route to the AM console with these parameters:

- Hostname: login.myproject.iam.example.com
- Path: /
- Target port: 80
- Secure route: (Selected)
- TLS Termination: Edge

b. Create a route to the IDM user interface with these parameters:

- Hostname: myproject.iam.example.com
- Path: openidm
- Target port: 80
- Secure route: (Selected)

- TLS Termination: Edge
- c. Create a route to the IDM console with these parameters:
- Hostname: myproject.iam.example.com
 - Path: /admin
 - Target port: 80
 - Secure route: (Selected)
 - TLS Termination: Edge

3.3. Verifying AM Access

Access the URL <https://login.myproject.iam.example.com/> in your browser to verify that you can access AM user interface.

If you are redirected to <http://openam>, then perform these steps:

1. Create a route for <http://openam/> in Minishift using the steps in the "To Create Routes" procedure.
2. Access the AM administration console, and update the Fixed value base URL parameter value to <https://myproject.iam.example.com>.
3. Save the changes and restart the AM pod using the **kubect**l command:

```
$ kubectl delete pod openam-pod
```

3.4. Deploying Identity Management

The example IDM deployment in Minishift uses Postgres to store data. Because of this, you'll need to deploy the `postgres-openidm` Helm chart before you deploy the `openidm` Helm chart.

To Deploy IDM

1. Deploy Postgres:

```
$ helm template postgres-openidm | oc apply -f -
```

2. Deploy IDM:

```
$ helm template openidm | oc apply -f -
```

3.5. Verifying IDM Access

Access the URL <https://myproject.iam.example.com/admin#login/> in your browser to verify that you can access the IDM console.

Appendix A. Getting Support

This appendix contains information about support options for the ForgeRock DevOps Examples and the ForgeRock Identity Platform.

A.1. ForgeRock DevOps Support

ForgeRock has developed artifacts in the `forgeops` and `forgeops-init` Git repositories for the purpose of deploying the ForgeRock Identity Platform in the cloud. The companion ForgeRock DevOps documentation provides examples, including the ForgeRock Cloud Deployment Model (CDM), to help you get started.

These artifacts and documentation are provided on an "as is" basis. ForgeRock does not guarantee the individual success developers may have in implementing the code on their development platforms or in production configurations.

A.1.1. Commercial Support

ForgeRock provides commercial support for the following DevOps resources:

- Dockerfiles and Helm charts in the `forgeops` Git repository
- ForgeRock [DevOps guides](#).

ForgeRock provides commercial support for the ForgeRock Identity Platform. For supported components, containers, and Java versions, see the following:

- *ForgeRock Access Management Release Notes*
- *ForgeRock Identity Management Release Notes*

- [ForgeRock Directory Services Release Notes](#)
- [ForgeRock Identity Message Broker Release Notes](#)
- [ForgeRock Identity Gateway Release Notes](#)

A.1.2. Support Limitations

ForgeRock provides no commercial support for the following:

- Artifacts other than Dockerfiles or Helm charts in the `forgeops` and `forgeops-init` repositories. Examples include scripts, example configurations, and so forth.
- Non-ForgeRock infrastructure. Examples include Docker, Kubernetes, Google Cloud Platform, Amazon Web Services, and so forth.
- Non-ForgeRock software. Examples include Java, Apache Tomcat, NGINX, Apache HTTP Server, and so forth.
- Production deployments that use the DevOps evaluation-only Docker images. When deploying the ForgeRock Identity Platform using Docker images, you must build and use your own images for production deployments. For information about how to build Docker images for the ForgeRock Identity Platform, see "[Building and Pushing Docker Images](#)" in the *DevOps Developer's Guide*.

A.1.3. Third-Party Kubernetes Services

ForgeRock supports deployments on Google Kubernetes Engine (GKE), Amazon Elastic Kubernetes Service (Amazon EKS), Microsoft Azure Kubernetes Service (AKS), and Red Hat OpenShift.

Red Hat OpenShift is a tested and supported platform using Kubernetes for deployment. ForgeRock uses OpenShift tools such as Minishift, as well as other representative environments such as Amazon AWS for the testing. We do not test using bare metal due to the many customer permutations of deployment and configuration that may exist, and therefore cannot guarantee that we have tested in the same way a customer chooses to deploy. We will make commercially reasonable efforts to provide first-line support for any reported issue. In the case we are unable to reproduce a reported issue internally, we will request the customer engage OpenShift support to collaborate on problem identification and remediation. Customers deploying on OpenShift are expected to have a support contract in place with IBM/Red Hat that ensures support resources can be engaged if this situation may occur.

A.2. Accessing Documentation Online

ForgeRock publishes comprehensive documentation online:

- The ForgeRock [Knowledge Base](#) offers a large and increasing number of up-to-date, practical articles that help you deploy and manage ForgeRock software.

While many articles are visible to community members, ForgeRock customers have access to much more, including advanced information for customers using ForgeRock software in a mission-critical capacity.

- ForgeRock product documentation, such as this document, aims to be technically accurate and complete with respect to the software documented. It is visible to everyone and covers all product features and examples of how to use them.

A.3. How to Report Problems or Provide Feedback

If you are a named customer Support Contact, contact ForgeRock using the [Customer Support Portal](#) to request information or report a problem with Dockerfiles or Helm charts in the DevOps Examples or the CDM.

If you have questions regarding the DevOps Examples or the CDM that are not answered in the documentation, file an issue at <https://github.com/ForgeRock/forgeops/issues>.

When requesting help with a problem, include the following information:

- Description of the problem, including when the problem occurs and its impact on your operation.
- Steps to reproduce the problem.

If the problem occurs on a Kubernetes system other than Minikube, GKE, EKS, OpenShift, or AKS, we might ask you to reproduce the problem on one of those.

- HTML output from the **debug-logs.sh** script. For more information, see "Running the debug-logs.sh Script" in the *DevOps Developer's Guide*.
- Description of the environment, including the following information:
 - Environment type: Minikube, GKE, EKS, AKS, or OpenShift.
 - Software versions of supporting components:
 - Oracle VirtualBox (Minikube environments only).
 - Docker client (all environments).
 - Minikube (all environments).
 - **kubect**l command (all environments).
 - Kubernetes Helm (all environments).
 - Google Cloud SDK (GKE environments only).
 - Amazon AWS Command Line Interface (EKS environments only).

- Azure Command Line Interface (AKS environments only).
- **forgeops** repository branch.
- Any patches or other software that might be affecting the problem.

A.4. Getting Support and Contacting ForgeRock

ForgeRock provides support services, professional services, training through ForgeRock University, and partner services to assist you in setting up and maintaining your deployments. For a general overview of these services, see <https://www.forgerock.com>.

ForgeRock has staff members around the globe who support our international customers and partners. For details on ForgeRock's support offering, including support plans and service level agreements (SLAs), visit <https://www.forgerock.com/support>.

Glossary

affinity (AM)	<p>AM affinity based load balancing ensures that the CTS token creation load is spread over multiple server instances (the token origin servers). Once a CTS token is created and assigned to a session, all subsequent token operations are sent to the same token origin server from any AM node. This ensures that the load of CTS token management is spread across directory servers.</p> <p>Source: <i>Best practices for using Core Token Service (CTS) Affinity based load balancing in AM</i></p>
Amazon EKS	<p>Amazon Elastic Container Service for Kubernetes (Amazon EKS) is a managed service that makes it easy for you to run Kubernetes on Amazon Web Services without needing to set up or maintain your own Kubernetes control plane.</p> <p>Source: <i>What is Amazon EKS</i> in the Amazon EKS documentation.</p>
ARN (AWS)	<p>An Amazon Resource Name (ARN) uniquely identifies an Amazon Web Service (AWS) resource. AWS requires an ARN when you need to specify a resource unambiguously across all of AWS, such as in IAM policies and API calls.</p> <p>Source: <i>Amazon Resource Names (ARNs) and AWS Service Namespaces</i> in the AWS documentation.</p>
AWS IAM Authenticator for Kubernetes	<p>The AWS IAM Authenticator for Kubernetes is an authentication tool that enables you to use <i>Amazon Web Services (AWS)</i> credentials for authenticating to a Kubernetes cluster.</p> <p>Source: <i>AWS IAM Authenticator for Kubernetes</i> README file on GitHub.</p>

cloud-controller-manager	<p>The <code>cloud-controller-manager</code> daemon runs controllers that interact with the underlying cloud providers. <code>cloud-controller-manager</code> is an alpha feature introduced in Kubernetes release 1.6. The <code>cloud-controller-manager</code> daemon runs cloud-provider-specific controller loops only.</p> <p>Source: <i>cloud-controller-manager</i> section in the Kubernetes Concepts documentation.</p>
Cloud Developer's Kit (CDK)	<p>The developer artifacts in the <code>forgeops</code> Git repository, together with the ForgeRock Identity Platform documentation form the Cloud Developer's Kit (CDK). Use the CDK to stand up the platform in your developer environment.</p>
Cloud Deployment Model (CDM)	<p>The Cloud Deployment Model (CDM) is a common use ForgeRock Identity Platform architecture, designed to be easy to deploy and easy to replicate. The ForgeRock Cloud Deployment Team has developed Helm charts, Docker images, and other artifacts expressly to build the CDM.</p>
CloudFormation (AWS)	<p>CloudFormation is a service that helps you model and set up your Amazon Web Services (AWS) resources. You create a template that describes all the AWS resources that you want. AWS CloudFormation takes care of provisioning and configuring those resources for you.</p> <p>Source: <i>What is AWS CloudFormation?</i> in the AWS documentation.</p>
CloudFormation template (AWS)	<p>An AWS CloudFormation template describes the resources that you want to provision in your AWS stack. AWS CloudFormation templates are text files formatted in JSON or YAML.</p> <p>Source: <i>Working with AWS CloudFormation Templates</i> in the AWS documentation.</p>
cluster	<p>A container cluster is the foundation of Kubernetes Engine. A cluster consists of at least one <code>cluster master</code> and multiple worker machines called nodes. The Kubernetes objects that represent your containerized applications all run on top of a cluster.</p> <p>Source: <i>Container Cluster Architecture</i> in the Kubernetes Concepts documentation.</p>
cluster master	<p>A cluster master schedules, runs, scales and upgrades the workloads on all nodes of the cluster. The cluster master also manages network and storage resources for workloads.</p> <p>Source: <i>Container Cluster Architecture</i> in the Kubernetes Concepts documentation.</p>
ConfigMap	<p>A configuration map, called <code>ConfigMap</code> in Kubernetes manifests, binds the configuration files, command-line arguments, environment</p>

variables, port numbers, and other configuration artifacts to the assigned containers and system components at runtime. The configuration maps are useful for storing and sharing non-sensitive, unencrypted configuration information.

Source: *ConfigMap* in the [Kubernetes Concepts](#) documentation.

container

A container is an allocation of resources such as CPU, network I/O, bandwidth, block I/O, and memory that can be “contained” together and made available to specific processes without interference from the rest of the system.

Source *Container Cluster Architecture* in the [Google Cloud Platform](#) documentation

DaemonSet

A set of daemons, called **DaemonSet** in Kubernetes manifests, manages a group of replicated pods. Usually, the daemon set follows an one-pod-per-node model. As you add nodes to a node pool, the daemon set automatically distributes the pod workload to the new nodes as needed.

Source *DaemonSet* in the [Google Cloud Platform](#) documentation.

Deployment

A Kubernetes deployment represents a set of multiple, identical pods. A Kubernetes deployment runs multiple replicas of your application and automatically replaces any instances that fail or become unresponsive.

Source: *Deployment* in the [Google Cloud Platform](#) documentation.

deployment controller

A deployment controller provides declarative updates for pods and replica sets. You describe a desired state in a deployment object, and the deployment controller changes the actual state to the desired state at a controlled rate. You can define deployments to create new replica sets, or to remove existing deployments and adopt all their resources with new deployments.

Source: *Deployments* in the [Google Cloud Platform](#) documentation.

Docker Cloud

Docker Cloud provides a hosted registry service with build and testing facilities for Dockerized application images; tools to help you set up and manage host infrastructure; and application lifecycle features to automate deploying (and redeploying) services created from images.

Source: *About Docker Cloud* in the [Docker Cloud](#) documentation.

Docker container

A Docker container is a runtime instance of a [Docker image](#). A Docker container is isolated from other containers and its host machine. You can control how isolated your container’s network,

storage, or other underlying subsystems are from other containers or from the host machine.

Source: Containers section in the Docker architecture documentation.

Docker daemon

The Docker daemon (`dockerd`) listens for Docker API requests and manages Docker objects such as images, containers, networks, and volumes. A Docker daemon can also communicate with other Docker daemons to manage Docker services.

Source: *Docker daemon* section in the Docker Overview documentation.

Docker Engine

The Docker Engine is a client-server application with these components:

- A server, which is a type of long-running program called a daemon process (the `dockerd` command)
- A REST API, which specifies interfaces that programs can use to talk to the daemon and tell it what to do
- A command-line interface (CLI) client (the `docker` command)

Source: Docker Engine section in the Docker Overview documentation.

Dockerfile

A Dockerfile is a text file that contains the instructions for building a Docker image. Docker uses the Dockerfile to automate the process of building a Docker image.

Source: *Dockerfile* section in the Docker Overview documentation.

Docker Hub

Docker Hub provides a place for you and your team to build and ship Docker images. You can create public repositories that can be accessed by any other Docker Hub user, or you can create private repositories you can control access to.

An image is an application you would like to run. A container is a running instance of an image.

Source: *Overview of Docker Hub* section in the Docker Overview documentation.

Docker image

A Docker image is a read-only template with instructions for creating a Docker container. Often, an image is based on another image, with some additional customization.

A Docker image includes the application code, a runtime engine, libraries, environment variables, and configuration files that are required to run the application.

An image is an application you would like to run. A container is a running instance of an image.

Source: *Docker objects* section in the Docker Overview documentation. [Hello Whales: Images vs. Containers in Dockers](#).

Docker namespace

Docker namespaces provide a layer of isolation. When you run a container, Docker creates a set of namespaces for that container. Each aspect of a container runs in a separate namespace and its access is limited to that namespace.

The **PID** namespace is the mechanism for remapping process IDs inside the container. Other namespaces such as net, mnt, ipc, and uts provide the isolated environments we know as containers. The user namespace is the mechanism for remapping user IDs inside a container.

Source: *Namespaces* section in the Docker Overview documentation.

Docker registry

A Docker registry stores Docker images. Docker Hub and Docker Cloud are public registries that anyone can use, and Docker is configured to look for images on Docker Hub by default. You can also run your own private registry.

Source: *Docker registries* section in the Docker Overview documentation.

Docker repository

A Docker repository is a public, certified repository from vendors and contributors to Docker. It contains Docker images that you can use as the foundation to build your applications and services.

Source: *Repositories on Docker Hub* section in the Docker Overview documentation.

Docker service

In a distributed application, different pieces of the application are called “services.” Docker services are really just “containers in production.” A Docker service runs only one image, but it codifies the way that image runs including which ports to use, the number replicas the container should run, and so on. By default, the services are load-balanced across all worker nodes.

Source: *About services* in the Docker Get Started documentation.

dynamic volume provisioning

The process of creating storage volumes on demand is called dynamic volume provisioning. Dynamic volume provisioning allows storage

volumes to be created on-demand. It automatically provisions storage when it is requested by users.

Source: *Dynamic Volume Provisioning* in the Kubernetes Concepts documentation.

egress

An egress controls access to destinations outside the network from within a Kubernetes network. For an external destination to be accessed from a Kubernetes environment, the destination should be listed as an allowed destination in the whitelist configuration.

Source: *Network Policies* in the Kubernetes Concepts documentation.

firewall rule

A firewall rule lets you allow or deny traffic to and from your virtual machine instances based on a configuration you specify. Each Kubernetes network has a set of firewall rules controlling access to and from instances in its subnets. Each firewall rule is defined to apply to either incoming *glossary-ingress*(ingress) or outgoing (egress) traffic, not both.

Source: *Firewall Rules Overview* in the Google Cloud Platform documentation.

garbage collection

Garbage collection is the process of deleting unused objects. *Kubelets* perform garbage collection for containers every minute and garbage collection for images every five minutes. You can adjust the high and low threshold flags and garbage collection policy to tune image garbage collection.

Source: *Garbage Collection* in the Kubernetes Concepts documentation.

Google Kubernetes Engine (GKE)

The Google Kubernetes Engine (GKE) is an environment for deploying, managing, and scaling your containerized applications using Google infrastructure. The GKE environment consists of multiple machine instances grouped together to form a *container cluster*.

Source: *Kubernetes Engine Overview* in the Google Cloud Platform documentation.

ingress

An ingress is a collection of rules that allow inbound connections to reach the cluster services.

Source: *Ingress* in the Kubernetes Concepts documentation.

instance group

An instance group is a collection of instances of virtual machines. The instance groups enable you to easily monitor and control the group of virtual machines together.

	<p>Source: <i>Instance Groups</i> in the Google Cloud Platform documentation.</p>
instance template	<p>An instance template is a global API resource that you can use to create VM instances and managed instance groups. Instance templates define the machine type, image, zone, labels, and other instance properties. They are very helpful in replicating the environments.</p> <p>Source: <i>Instance Templates</i> in the Google Cloud Platform documentation.</p>
kubectrl	<p>The kubectrl command-line tool supports several different ways to create and manage Kubernetes objects.</p> <p>Source: <i>Kubernetes Object Management</i> in the Kubernetes Concepts documentation.</p>
kube-controller-manager	<p>The Kubernetes controller manager is a process that embeds core controllers that are shipped with Kubernetes. Logically each controller is a separate process, but to reduce complexity, they are all compiled into a single binary and run in a single process.</p> <p>Source: <i>kube-controller-manager</i> in the Kubernetes Reference documentation.</p>
kubelet	<p>A kubelet is an agent that runs on each node in the cluster. It ensures that containers are running in a pod.</p> <p>Source: <i>kubelets</i> in the Kubernetes Concepts documentation.</p>
kube-scheduler	<p>The kube-scheduler component is on the master node and watches for newly created pods that do not have a node assigned to them, and selects a node for them to run on.</p> <p>Source: <i>Kubernetes components</i> in the Kubernetes Concepts documentation.</p>
Kubernetes	<p>Kubernetes is an open source platform designed to automate deploying, scaling, and operating application containers.</p> <p>Source: <i>Kubernetes Concepts</i></p>
Kubernetes DNS	<p>A Kubernetes DNS pod is a pod used by the kubelets and the individual containers to resolve DNS names in the cluster.</p> <p>Source: <i>DNS for services and pods</i> in the Kubernetes Concepts documentation.</p>

Kubernetes namespace	<p>A Kubernetes namespace is a virtual cluster that provides a way to divide cluster resources between multiple users. Kubernetes starts with three initial namespaces:</p> <ul style="list-style-type: none">• default: The default namespace for user created objects which don't have a namespace• kube-system: The namespace for objects created by the Kubernetes system• kube-public: The automatically created namespace that is readable by all users <p>Kubernetes supports multiple virtual clusters backed by the same physical cluster.</p> <p>Source: <i>Namespaces</i> in the Kubernetes Concepts documentation.</p>
Let's Encrypt	<p>Let's Encrypt is a free, automated, and open certificate authority.</p> <p>Source: Let's Encrypt web site.</p>
network policy	<p>A Kubernetes network policy specifies how groups of pods are allowed to communicate with each other and with other network endpoints.</p> <p>Source: <i>Network policies</i> in the Kubernetes Concepts documentation.</p>
node (Kubernetes)	<p>A Kubernetes node is a virtual or physical machine in the cluster. Each node is managed by the master components and includes the services needed to run the pods.</p> <p>Source: <i>Nodes</i> in the Kubernetes Concepts documentation.</p>
node controller (Kubernetes)	<p>A Kubernetes node controller is a Kubernetes master component that manages various aspects of the nodes such as: lifecycle operations on the nodes, operational status of the nodes, and maintaining an internal list of nodes.</p> <p>Source: <i>Node Controller</i> in the Kubernetes Concepts documentation.</p>
persistent volume	<p>A persistent volume (PV) is a piece of storage in the cluster that has been provisioned by an administrator. It is a resource in the cluster just like a node is a cluster resource. PVs are volume plugins that have a lifecycle independent of any individual pod that uses the PV.</p> <p>Source: <i>Persistent Volumes</i> in the Kubernetes Concepts documentation.</p>
persistent volume claim	<p>A persistent volume claim (PVC) is a request for storage by a user. A PVC specifies size, and access modes such as:</p>

- Mounted once for read and write access
- Mounted many times for read-only access

Source: *Persistent Volumes* in the Kubernetes Concepts documentation.

pod anti-affinity
(Kubernetes)

Kubernetes pod anti-affinity allows you to constrain which nodes can run your pod, based on labels on the **Pods** that are already running on the node rather than based on labels on nodes. Pod anti-affinity enables you to control the spread of workload across nodes and also isolate failures to nodes.

Source: *Inter-pod affinity and anti-affinity*

pod (Kubernetes)

A Kubernetes pod is the smallest, most basic deployable object in Kubernetes. A pod represents a single instance of a running process in a cluster. Containers within a pod share an IP address and port space.

Source: *Understanding Pods* in the Kubernetes Concepts documentation.

replication controller

A replication controller ensures that a specified number of Kubernetes pod replicas are running at any one time. The **replication controller** ensures that a pod or a homogeneous set of pods is always up and available.

Source: *ReplicationController* in the Kubernetes Concepts documentation.

secret (Kubernetes)

A Kubernetes secret is a secure object that stores sensitive data, such as passwords, OAuth 2.0 tokens, and SSH keys in your clusters.

Source: *Secrets* in the Kubernetes Concepts documentation.

security group (AWS)

A security group acts as a virtual firewall that controls the traffic for one or more compute instances.

Source: *Amazon EC2 Security Groups* in the AWS documentation.

service (Kubernetes)

A Kubernetes service is an abstraction which defines a logical set of pods and a policy by which to access them. This is sometimes called a microservice.

Source: *Services* in the Kubernetes Concepts documentation.

shard

Sharding is a way of partitioning directory data so that the load can be shared by multiple directory servers. Each data partition, also

known as a *shard*, exposes the same set of naming contexts, but only a subset of the data. For example, a distribution might have two shards. The first shard contains all users whose name begins with A-M, and the second contains all users whose name begins with N-Z. Both have the same naming context.

Source: *Class Partition* in the *OpenDJ Javadoc*.

stack (AWS)

A stack is a collection of AWS resources that you can manage as a single unit. You can create, update, or delete a collection of resources by using stacks. All the resources in a stack are defined by the [template](#).

Source: *Working with Stacks* in the AWS documentation.

stack set (AWS)

A stack set is a container for stacks. You can provision stacks across AWS accounts and regions by using a single AWS [template](#). All the resources included in each stack of a stack set are defined by the same template.

Source: *StackSets Concepts* in the AWS documentation.

volume (Kubernetes)

A Kubernetes volume is a storage volume that has the same lifetime as the pod that encloses it. Consequently, a volume outlives any containers that run within the pod, and data is preserved across container restarts. When a pod ceases to exist, the Kubernetes volume also ceases to exist.

Source: *Volumes* in the Kubernetes Concepts documentation.

VPC (AWS)

A virtual private cloud (VPC) is a virtual network dedicated to your AWS account. It is logically isolated from other virtual networks in the AWS Cloud.

Source: *What Is Amazon VPC?* in the AWS documentation.

worker node (AWS)

An Amazon Elastic Container Service for Kubernetes (Amazon EKS) worker node is a standard compute instance provisioned in Amazon EKS.

Source: *Worker Nodes* in the AWS documentation.

workload (Kubernetes)

A Kubernetes workload is the collection of applications and batch jobs packaged into a [container](#). Before you deploy a workload on a cluster, you must first package the workload into a container.

Source: *Understanding Pods* in the Kubernetes Concepts documentation.